

P&S Modern SSDs

Advanced NAND Flash Commands & Address Translation

Dr. Jisung Park

Prof. Onur Mutlu

ETH Zürich

Spring 2022

4 April 2022

Recap: SSD & NAND Flash Memory

- SSD organization
 - SSD controller: Multicore CPU + per-channel flash controllers
 - DRAM: Metadata store, 0.1% of SSD capacity
 - NAND flash chips
 - Channel (Package(s)) – Die (Chip) – Plane – Block – Page
- NAND flash characteristics
 - Erase-before-write, asymmetry in operation units (read/program: page, erase: block), limited endurance, retention loss...
- Basic NAND flash operations
 - Read/program/erase

Today's Agenda

- Advanced NAND Flash Commands
- Address Translation & Garbage Collection

SSD Performance

■ Latency (or response time)

- The time delay **until the request is returned**

- Average read latency (4 KiB): **67 us**

- Average write latency (4 KiB): **47 us**

HDD:
5~8 ms

■ Throughput

- The **number of requests** that can be serviced per unit time

 - **IOPS:** Input/output Operations Per Second

- **Random read** throughput: up to **500K IOPS**

- **Random write** throughput: up to **480K IOPS**

HDD:
> 1K IOPS

■ Bandwidth

- The **amount of data** that can be accessed per unit time

- **Sequential read** bandwidth: up to **3,500 MB/s**

- **Sequential write** bandwidth: up to **3,000 MB/s**

HDD:
~100 MB/s

Source: <https://www.anandtech.com/show/16504/the-samsung-ssd-980-500gb-1tb-review>

NAND Flash Chip Performance

■ Chip operation latency

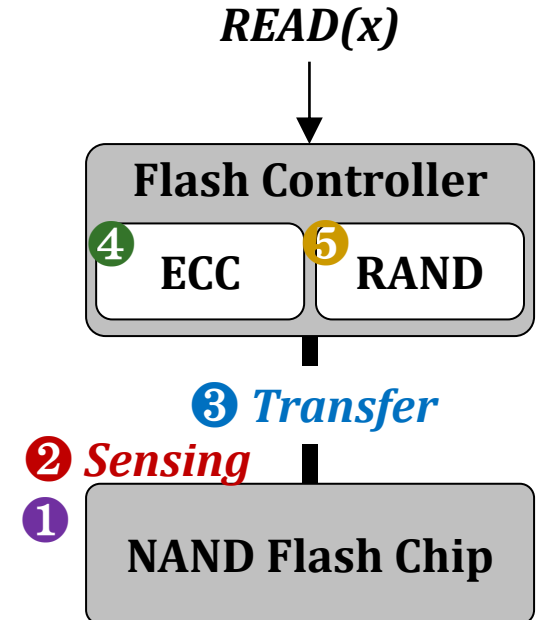
- ❑ **tR**: Latency of reading (sensing) data from the cells **into the on-chip page buffer**
- ❑ **tPROG**: Latency of programming the cells **with data in the page buffer**
- ❑ **tBERS**: Latency of erasing the cells (block)
- ❑ Varies depending on the **MLC technology, processing node, and microarchitecture**
 - In 3D TLC NAND flash, $tR/tPROG/tBERS \approx 100\mu s/700\mu s/3ms$

■ I/O rate

- ❑ **Number of bits** transferred via **a single I/O pin** per unit time
- ❑ A typical flash chip transfers data in **a byte granularity** (i.e., via 8 I/O pins)
- ❑ e.g., 1-Gb I/O rate & 16-KiB page size $\rightarrow tDMA = 16 \mu s$

NAND Flash Chip Performance (Cont.)


- t_R , t_{PROG} , and t_{BERS}
 - Latencies for chip-level read/program/erase operations
 - t_R : 50~100 us
 - t_{PROG} : 700us~1000 us
 - t_{BERS} : 3ms~5ms
- Flash-controller level latency
 - 1-Gb I/O rate and 16-KiB page size
 - Read
 - $(t_{CMD}) + t_R + t_{DMA} + t_{ECC_{DEC}} + (t_{RND})$
 - e.g., $100 + 16 + 20 = 136$ us



NAND Flash Chip Performance (Cont.)

- t_R , t_{PROG} , and t_{BERS}
 - Latencies for chip-level read/program/erase operations
 - t_R : 50~100 μs
 - t_{PROG} : 700 μs ~1000 μs
 - t_{BERS} : 3ms~5ms

- Flash-controller level latency
 - 1-Gb I/O rate and 16-KiB page size
 - Read
 - $(t_{\text{CMD}}) + t_R + t_{\text{DMA}} + t_{\text{ECC}_{\text{DEC}}} + (t_{\text{RND}})$
 - e.g., $100 + 16 + 20 = 136 \mu\text{s}$
 - Program
 - $(t_{\text{RND}}) + t_{\text{ECC}_{\text{ENC}}} + (t_{\text{CMD}}) + t_{\text{DMA}} + t_{\text{PROG}}$
 - e.g., $20 + 16 + 700 = 736 \mu\text{s}$



NAND Flash Chip Performance (Cont.)

■ How about bandwidth?

□ Read

■ 16 KiB / 136 μ s \approx 120 MB/s

□ Write

■ 16 KiB / 736 μ s \approx 22 MB/s

WAIT!

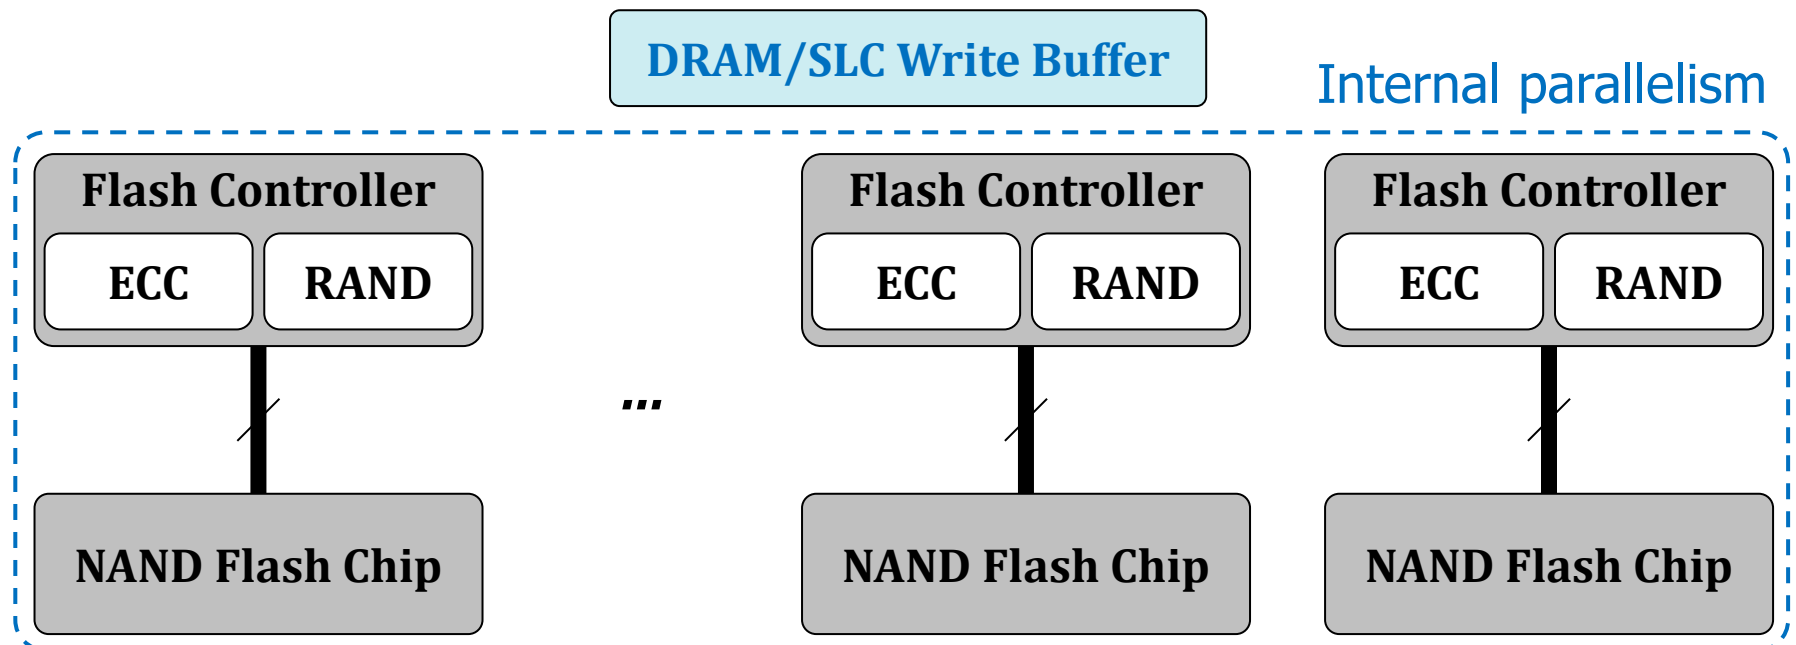
SSD read latency: 67 μ s

SSD read bandwidth: 3.5 GB/s

SSD write latency: 47 μ s

SSD write bandwidth: 3 GB/s

Optimizations w/ advanced commands




Advanced Commands for Small Reads

- Minimum I/O units in modern file systems: 4 KiB
 - Latency & bandwidth waste due to I/O-unit mismatch
 - e.g., A page read unnecessarily reads/transfers 12-KiB data
- Optimization 1: Sub-page sensing
 - e.g., Micron SNAP READ operation¹
 - Microarchitecture-level optimization – directly reduces tR
- Optimization 2: Random Data Out (RDO)
 - Data transfer with an arbitrary offset and size
 - Reduce tDMA and tECC_{DEC}


¹https://media-www.micron.com/-/media/client/global/documents/products/technical-note/nand-flash/tn_2993_snap_read.pdf

CACHE READ Command

- Performs consecutive reads in a pipelined manner



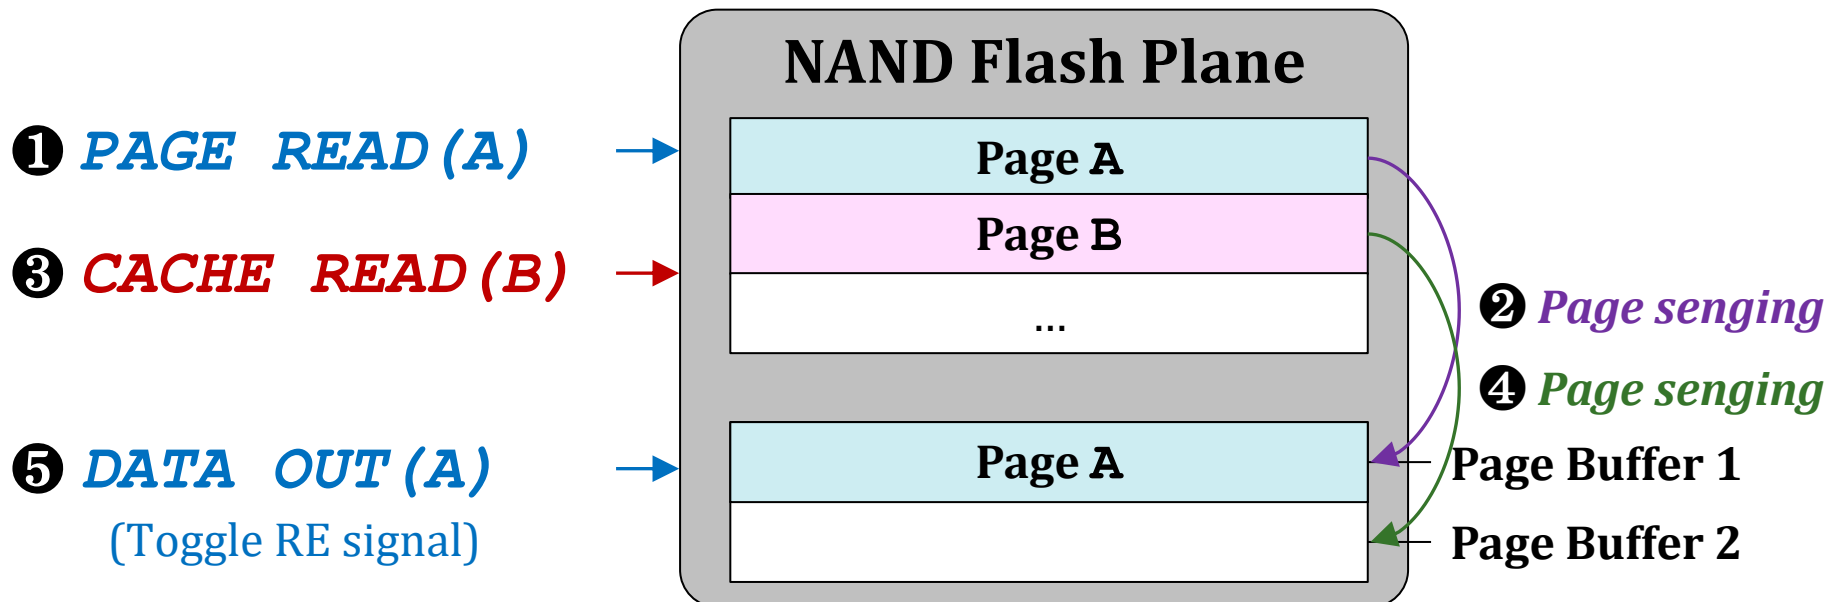

Regular PAGE READ:
Overlaps **only** t_{ECC} with t_R



CACHE READ:
Overlaps t_{DMA} & t_{ECC}
with t_R

Enabling the CACHE READ Command

- Needs additional on-chip page buffer




CACHE READ Command: Benefit

- Removes tDMA from the critical path
 - Increases throughput/bandwidth
 - Reduces effective latency
 - By reducing the time delay for a request being blocked by the previous request

Multi-Plane Operations

- Concurrent operations on different planes
 - Recall: Planes share WLs and row/column decoders



- Opportunity: Planes can **concurrently** operate
- Constraints: Only for **the same operations on the same page offset**

Multi-Plane Operations: Benefit

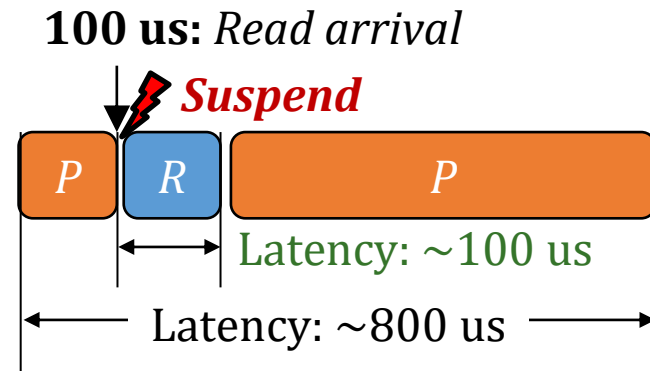

- Increase the throughput/bandwidth almost linearly with # of planes that concurrently operate
 - Bandwidth with regular page programs:
 $16 \text{ KiB} / 736 \text{ us} \approx 22 \text{ MB/s}$
 - Bandwidth with multi-plane page programs (2 plane):
 $32 \text{ KiB} / 736 + 16 (\text{tDMA}) + 20 (\text{tECC}) \text{ us} \approx 41.5 \text{ MB/s}$
- Per-operation latency increases
 - Regular page program: $\text{tECC}_{\text{ENC}} + \text{tDMA} + \text{tPROG}$
 - Multi-plane page program: $N_{\text{Plane}} \times (\text{tECC}_{\text{ENC}} + \text{tDMA}) + \text{tPROG}$
- The benefits highly depend on the access pattern and FTL's data placement
 - Random-read-dominant vs. Random-write-dominant

Program & Erase Suspensions

- Read performance is often more important
 - Writes can be done in an asynchronous manner using buffers
 - e.g., return a write request immediately after receiving the data (and storing it to the write buffer)
 - A read request can be returned only when the requested data is ready (after reading the data from the chip)
- Significant latency asymmetry
 - tR: 100 us, tPROG: 700 us, tBERS: 5 ms (TLC NAND flash)
 - If the chip is designed to program all the pages in the same WL at once, the actual program latency is 2,100 us
 - The worst-case chip-level read latency can be 50x longer than the best-case latency

Program & Erase Suspensions (Cont.)

- Suspends an on-going program (erase) operation once a read arrives



- Pros: Significantly decreases the read latency
- Cons
 - Additional page buffer (for data to program)
 - Complicated I/O scheduling (Until when can we suspend on-going program requests?)
 - Negative impact on the endurance

Summary

- **Subpage Sensing & Random Data Out (RDO)**
 - For **I/O-unit mismatch** b/w OS and NAND flash memory
- **Cache Read Command**
 - For improving **a chip's read throughput**
 - By overlapping data transfer and page sensing
- **Multi-Plane Operations**
 - For improving **a chip's throughput**
 - By enabling **concurrently operation of multiple planes**
- **Program & Erase Suspensions**
 - For improving **the read latency** (**operation latency asymmetry**)
 - By **prioritizing latency-sensitive reads** over writes/erases

Today's Agenda

- Advanced NAND Flash Commands
- Address Translation & Garbage Collection

Flash Translation Layer: Overview

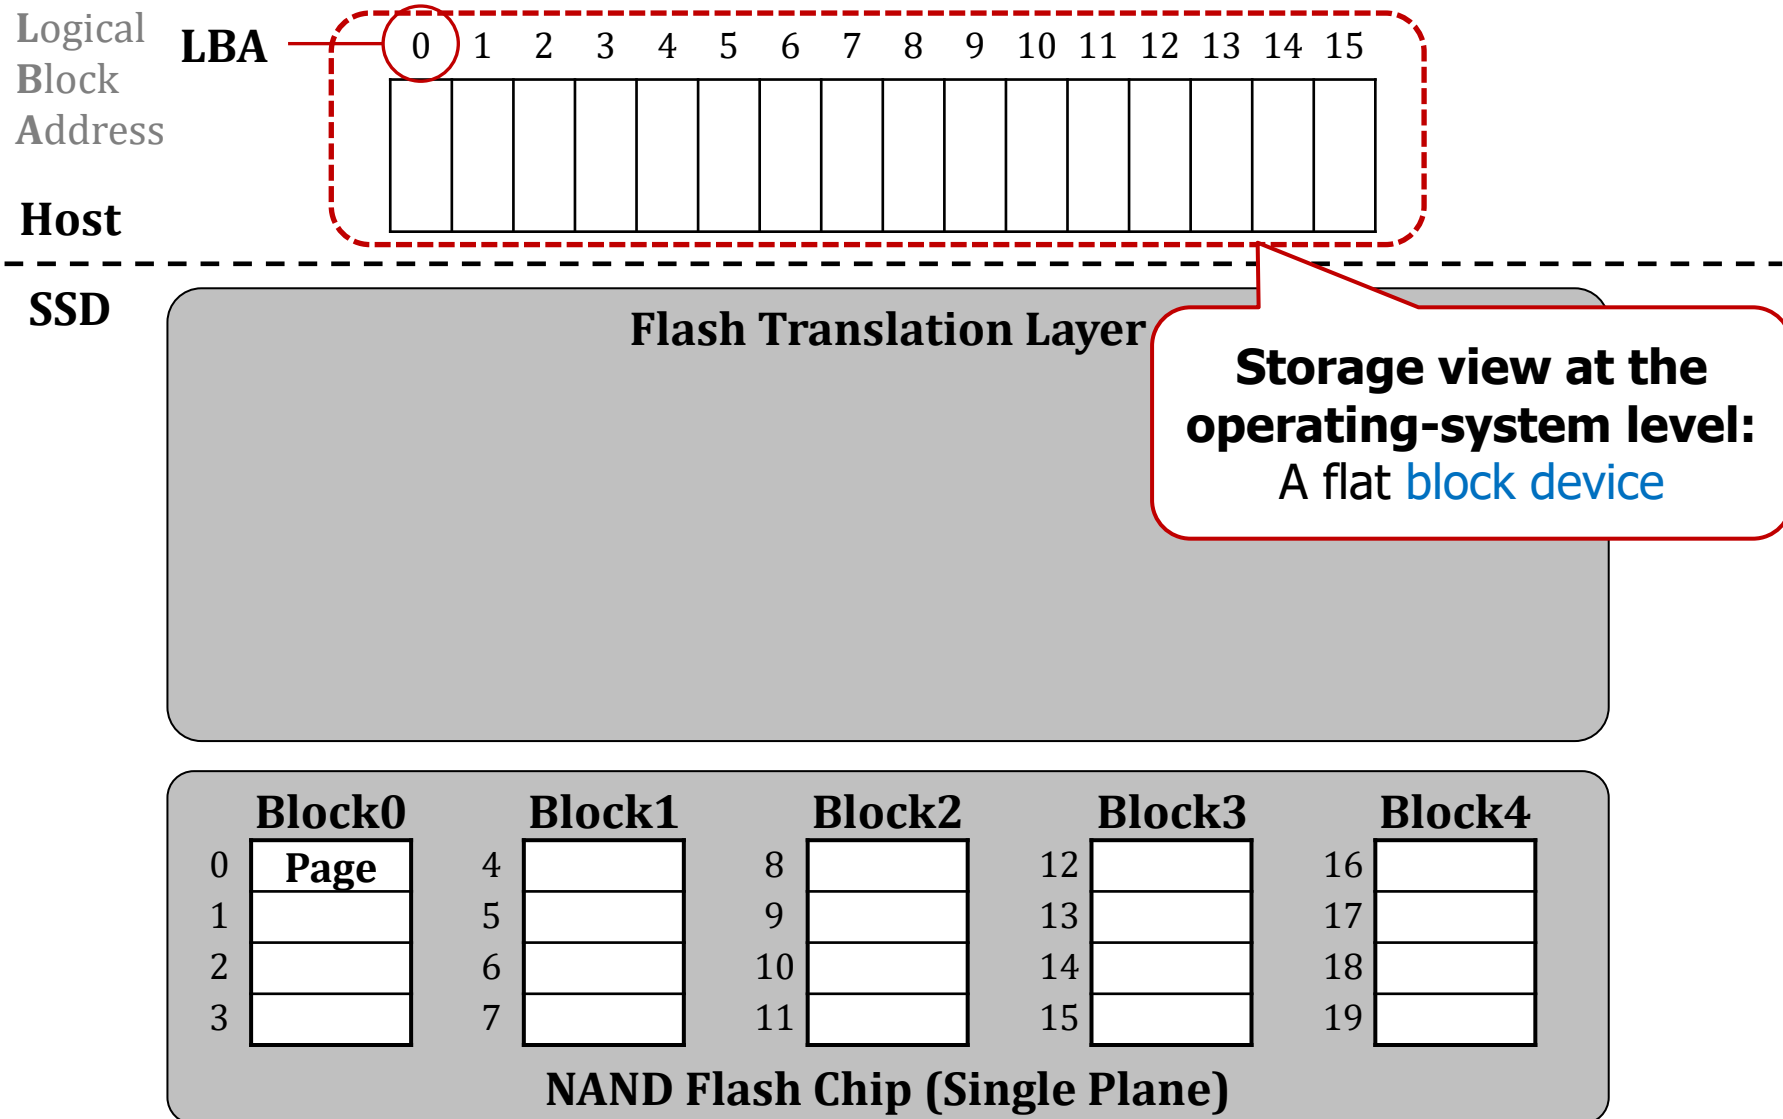
- SSD firmware (often referred to as SSD controller)
 - Provides **backward compatibility** with traditional HDDs
 - By **hiding unique characteristics** of NAND flash memory

- Responsible for many important **SSD-management tasks**
 - Address translation + garbage collection
 - Performs **out-of-place writes** due to erase-before-write property
 - Wear leveling
 - To prolong SSD lifetime by **evenly distributing** P/E cycles
 - Data refresh
 - Resets transient errors by **copying data** to a new page(s)
 - I/O scheduling
 - To take full advantage of **SSD internal parallelism**

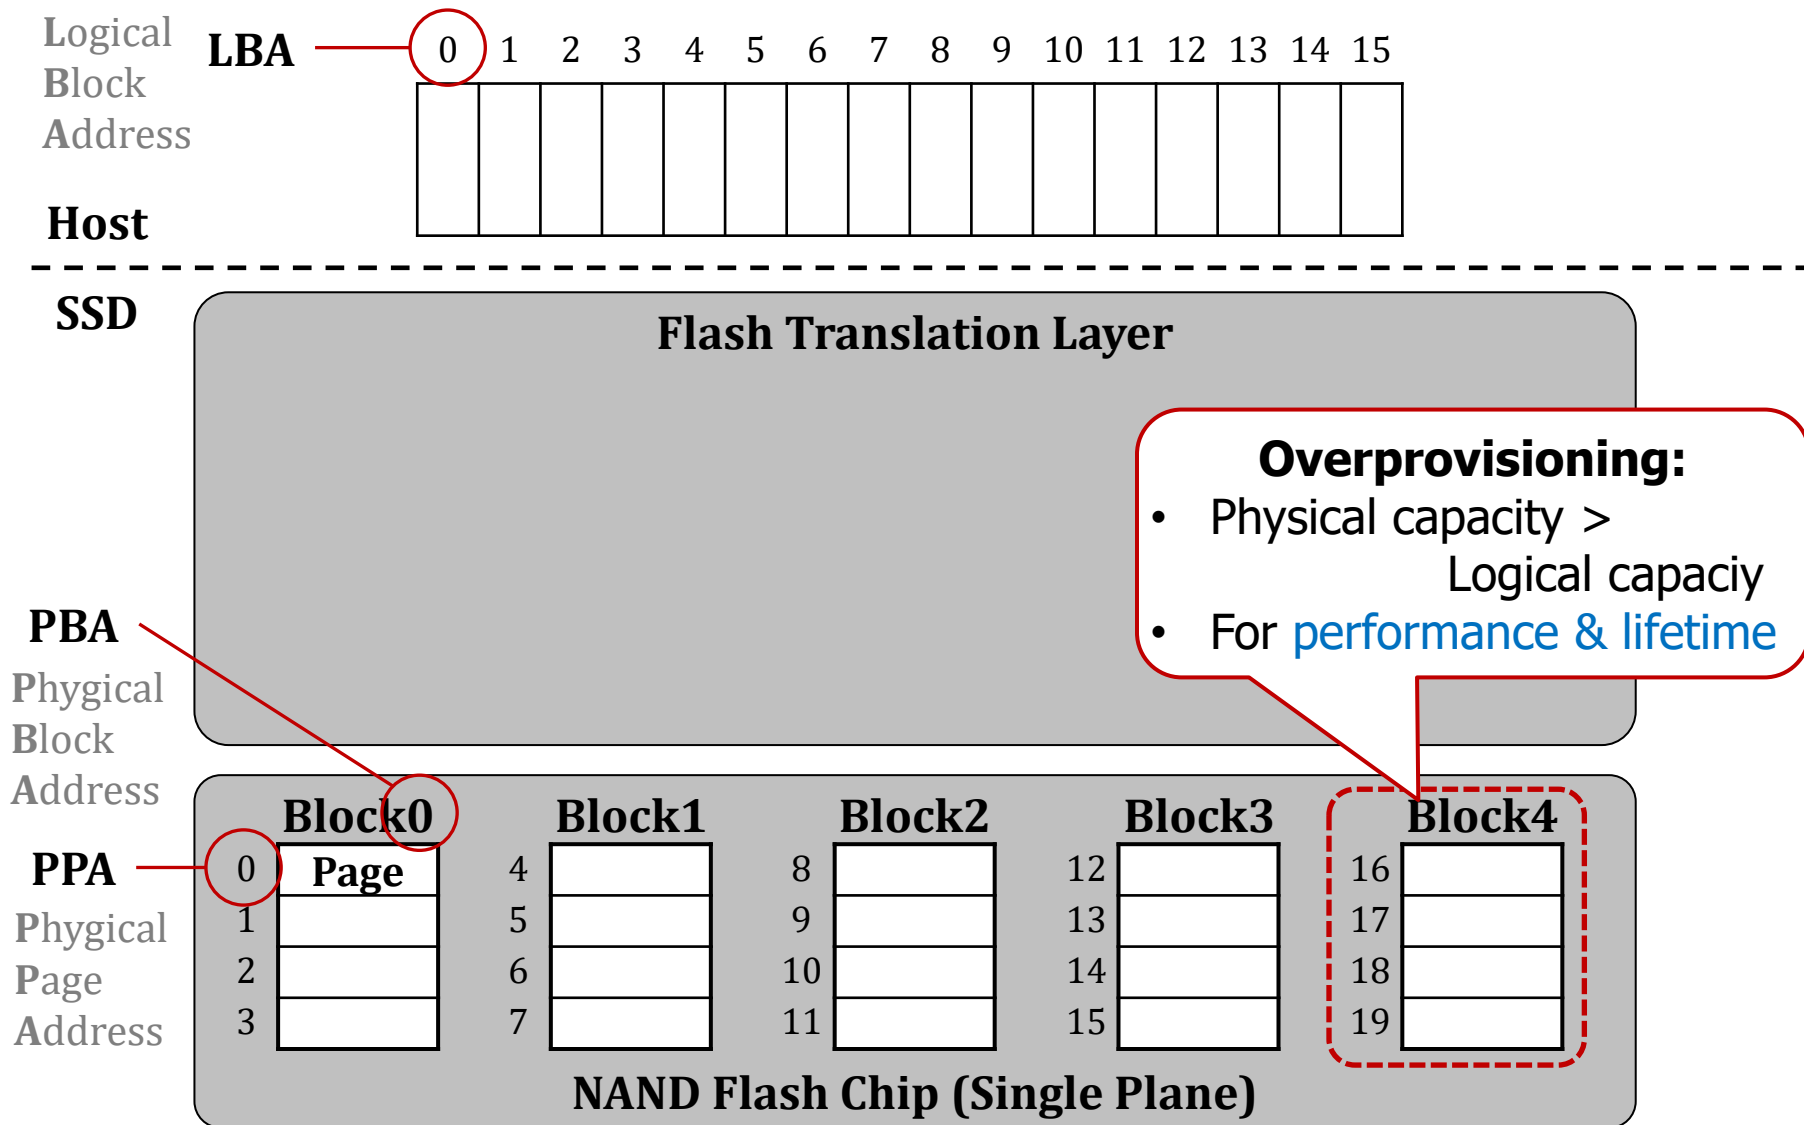
Flash Translation Layer: Overview

- SSD firmware (often referred to as SSD controller)
 - Provides **backward compatibility** with traditional HDDs
 - By **hiding unique characteristics** of NAND flash memory
- Responsible for many important **SSD-management tasks**
 - Address translation + garbage collection
 - Performs **out-of-place writes** due to erase-before-write property
 - Wear leveling
 - To prolong SSD lifetime by **evenly distributing** P/E cycles
 - Data refresh
 - Resets transient errors by **copying data** to a new page(s)
 - I/O scheduling
 - To take full advantage of **SSD internal parallelism**

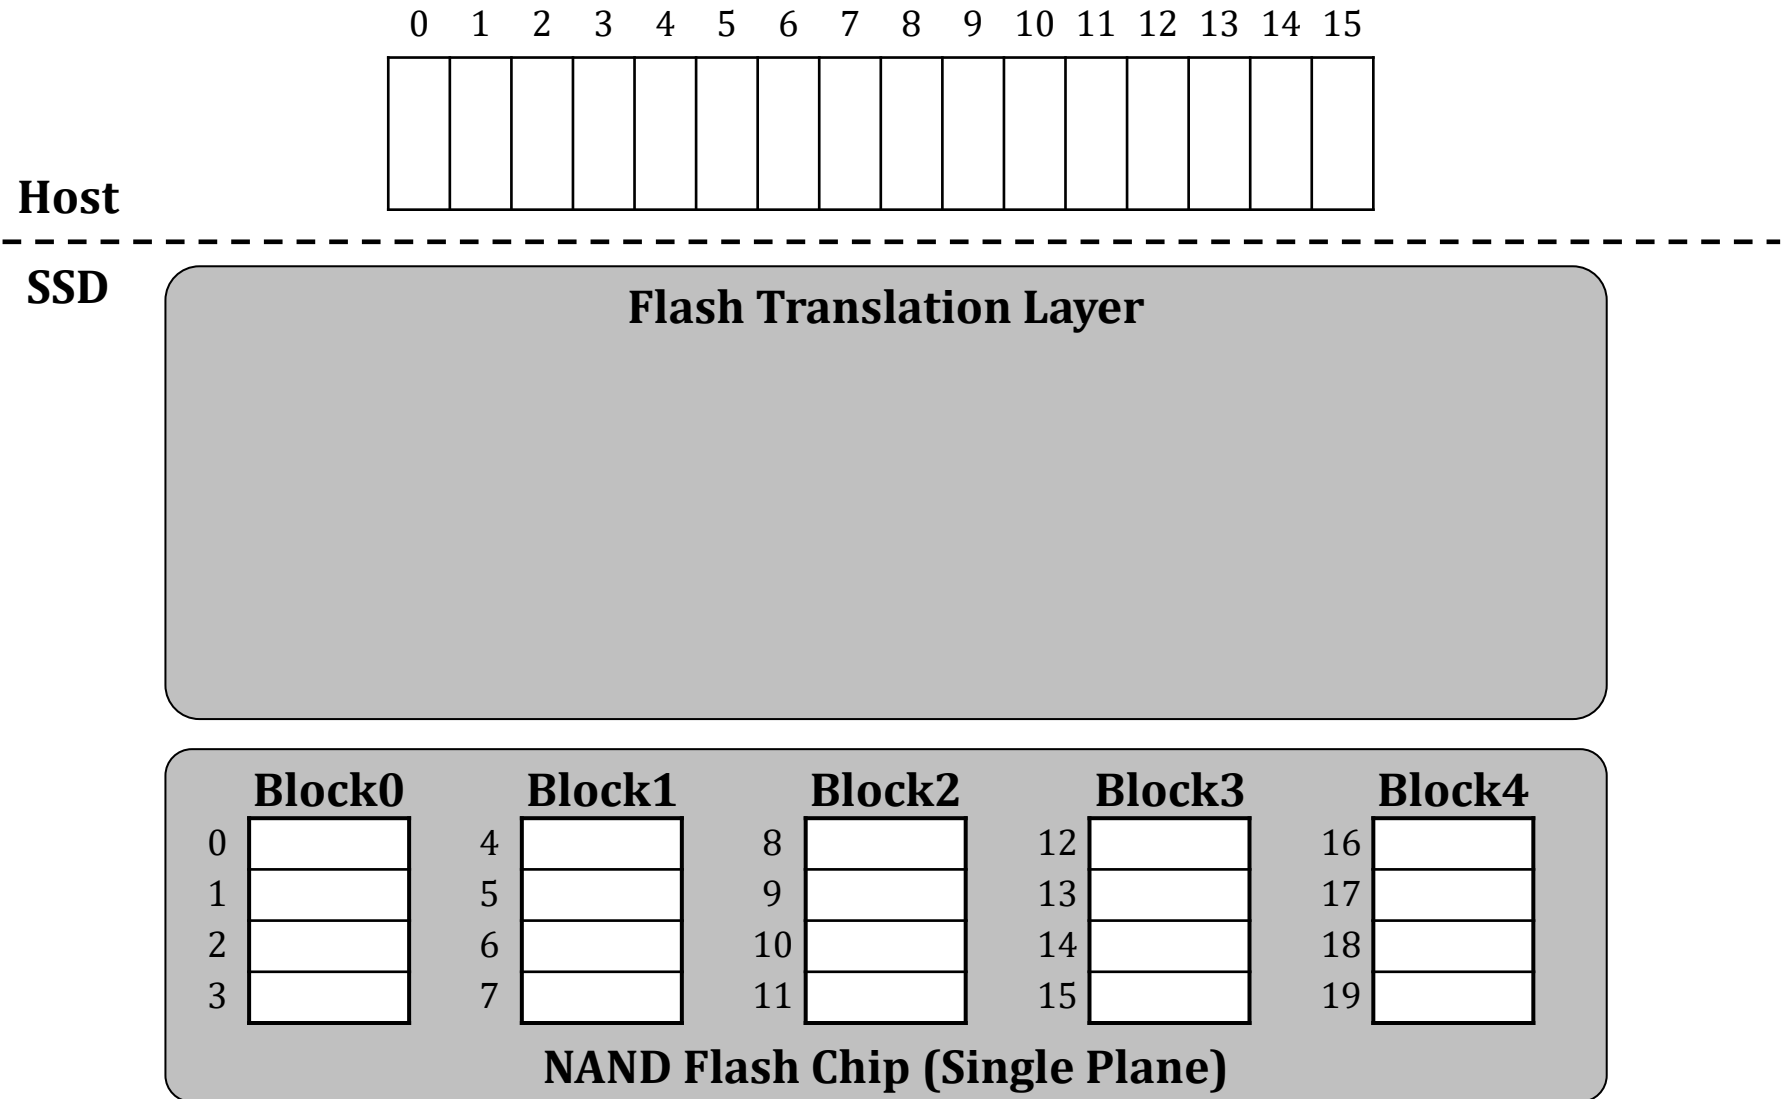
Simple SSD Architecture



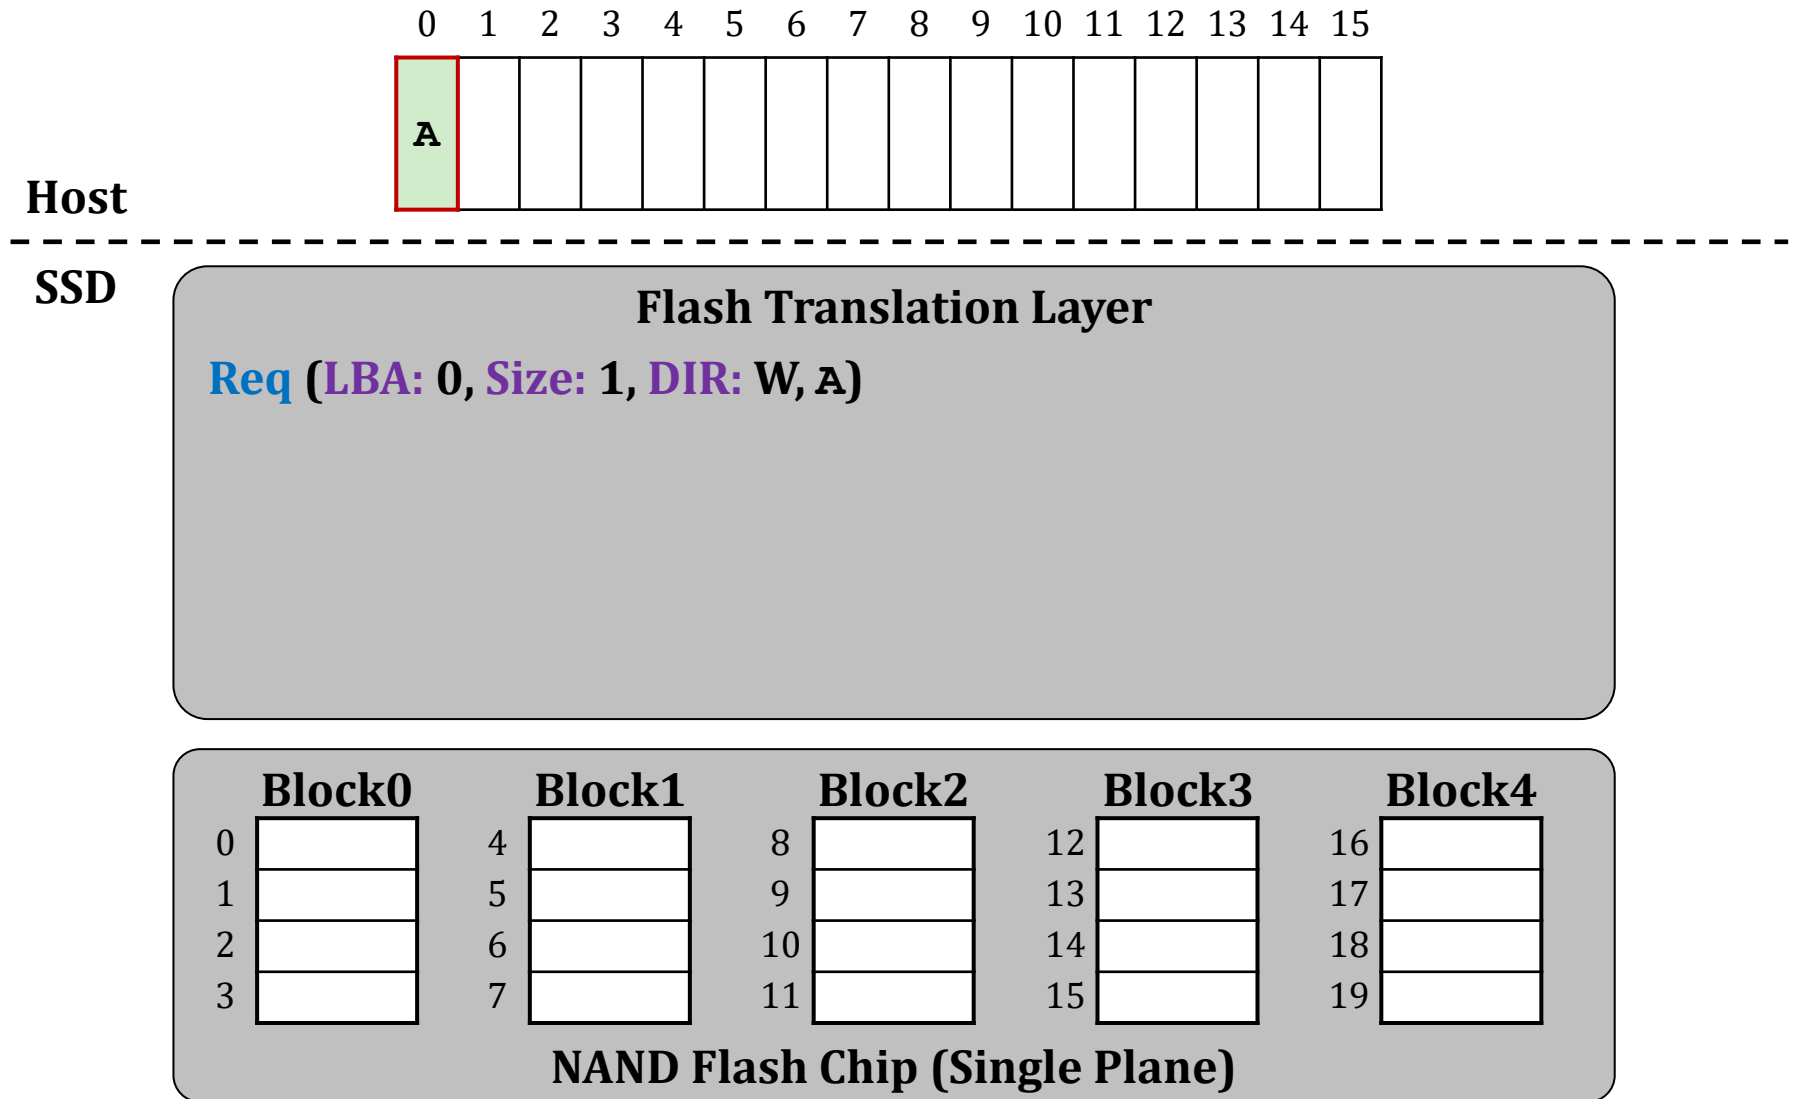
Simple SSD Architecture



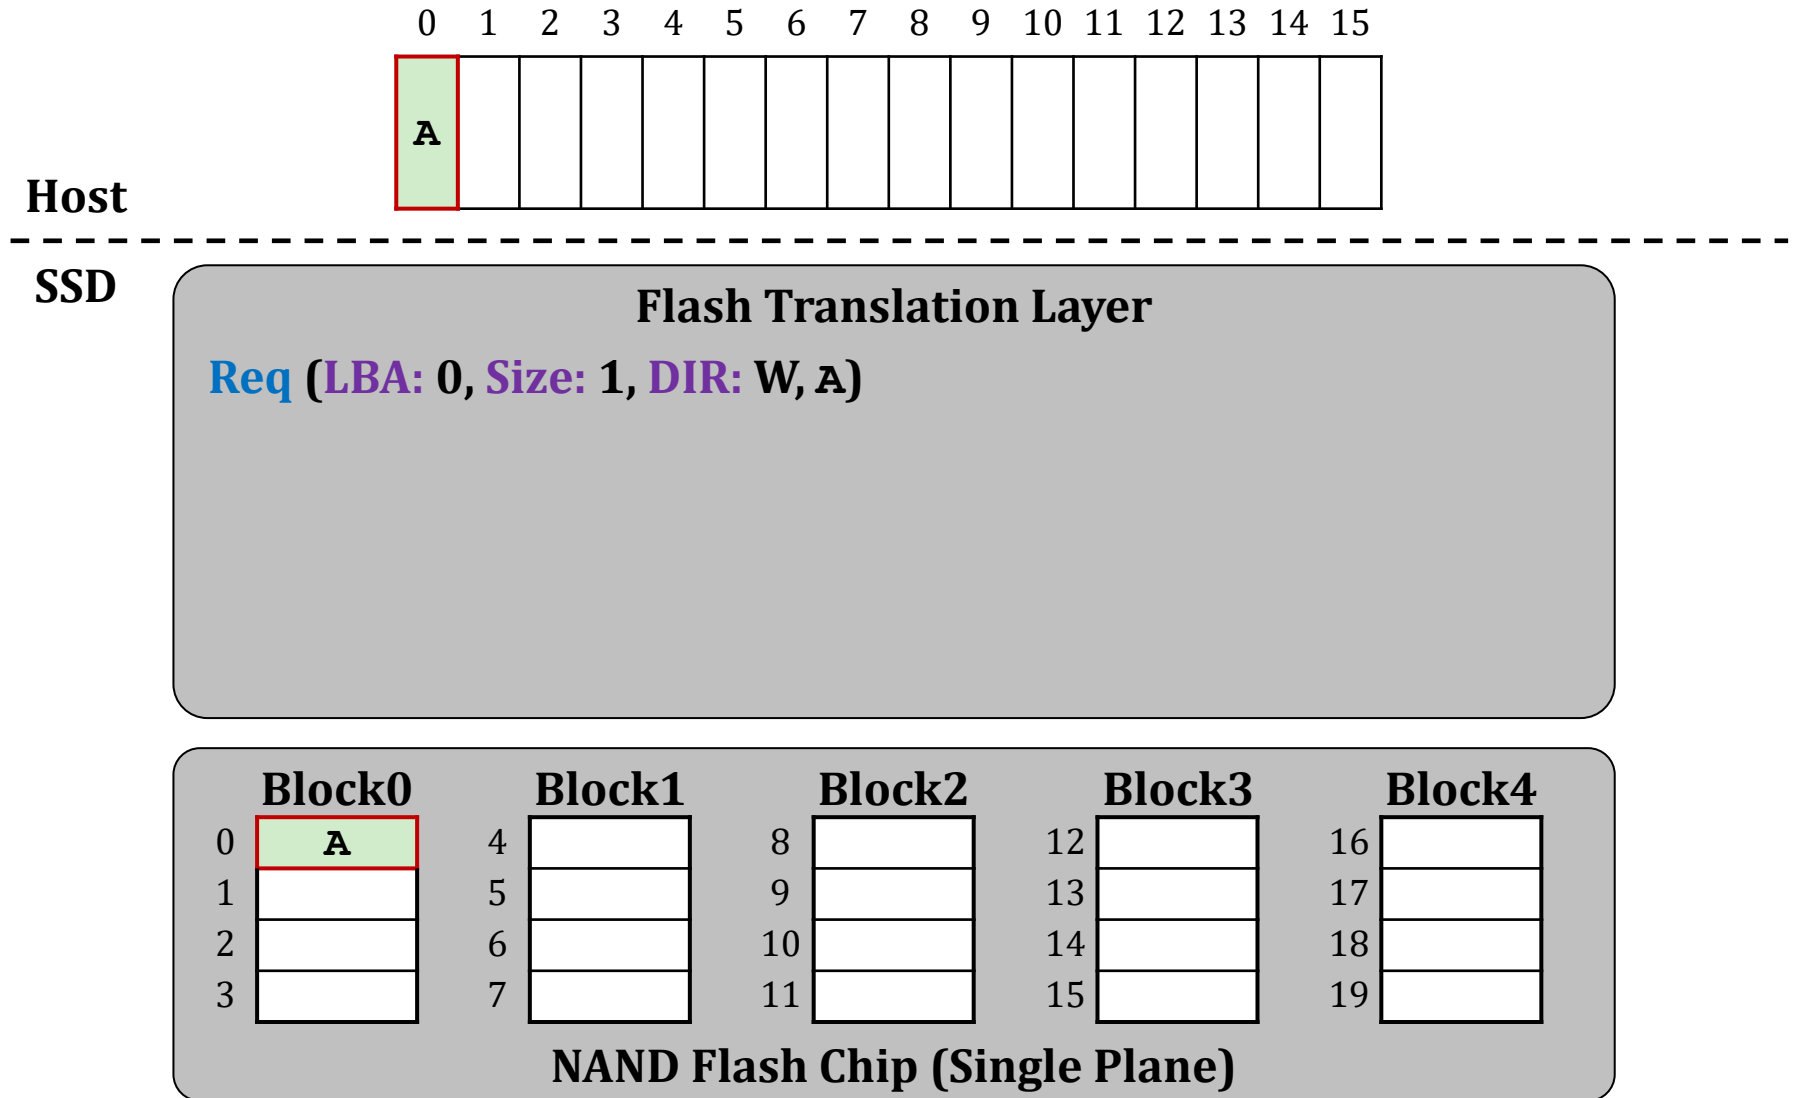
Write Request Handling: Page Write



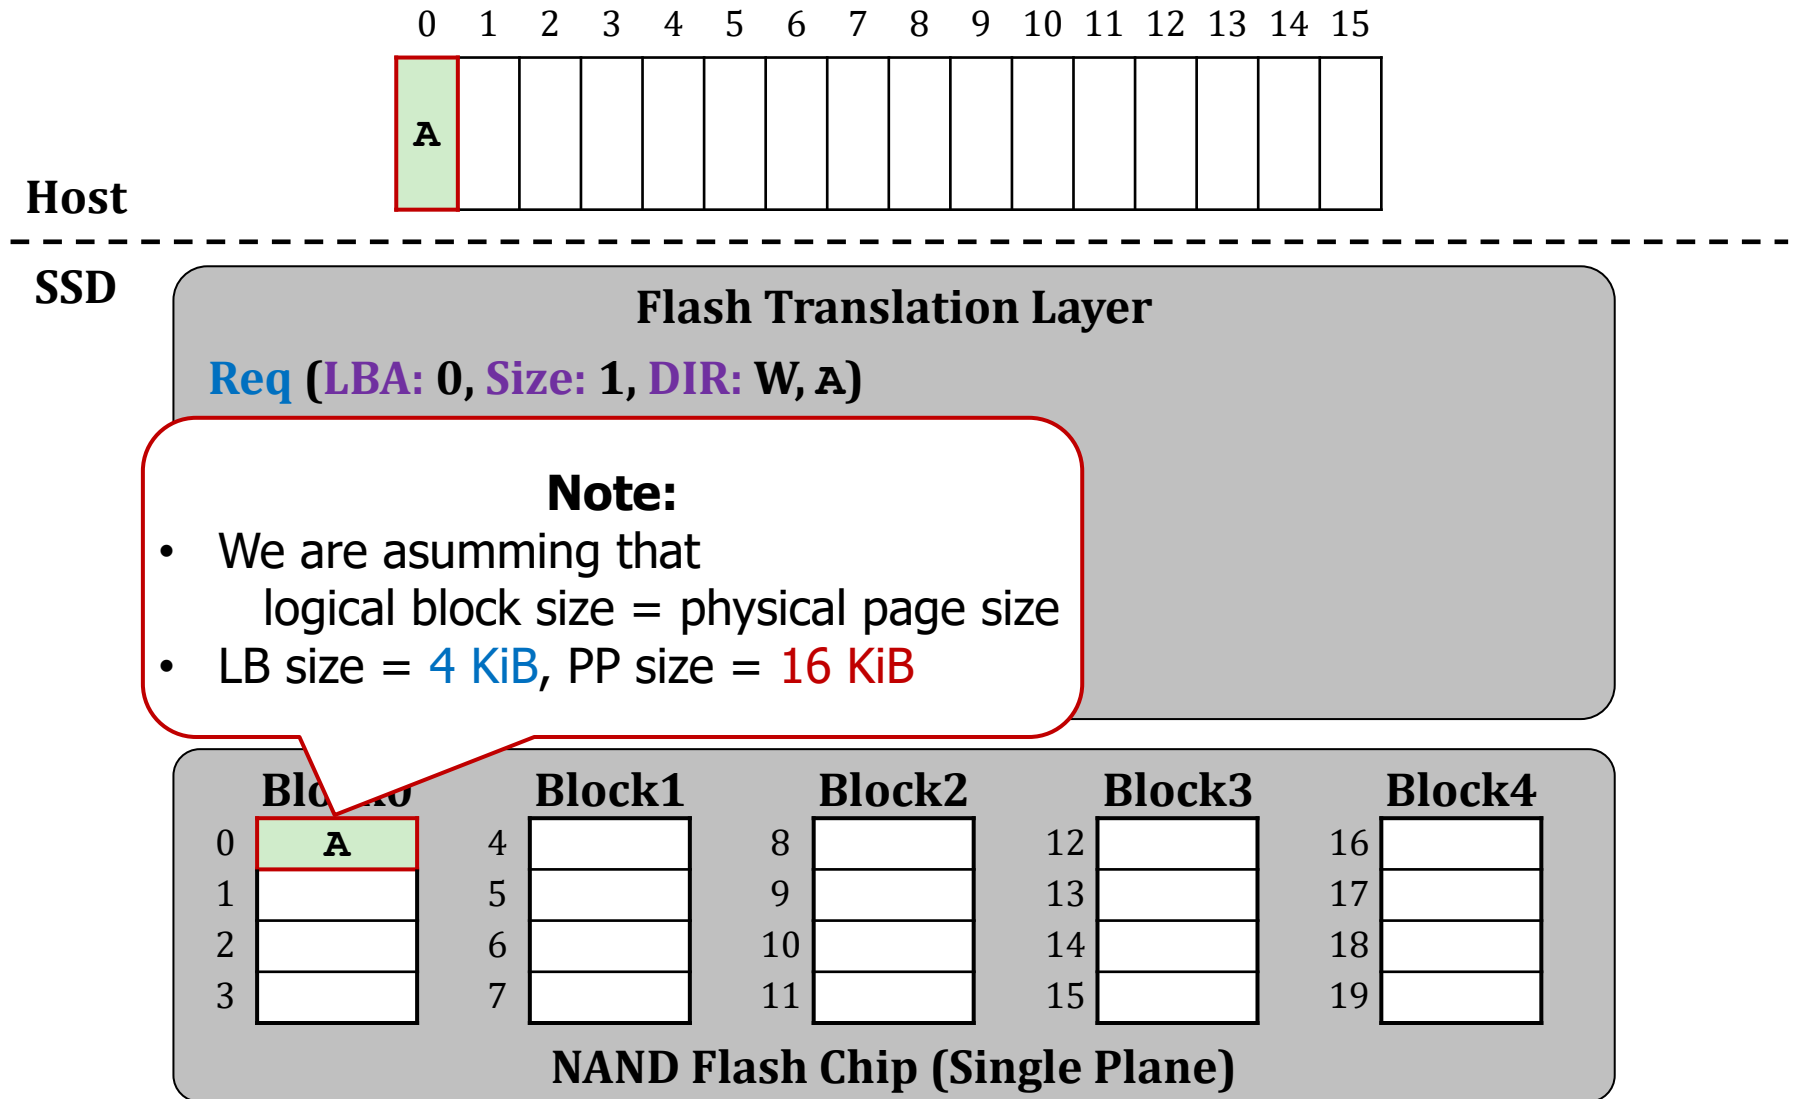
Write Request Handling: Page Write



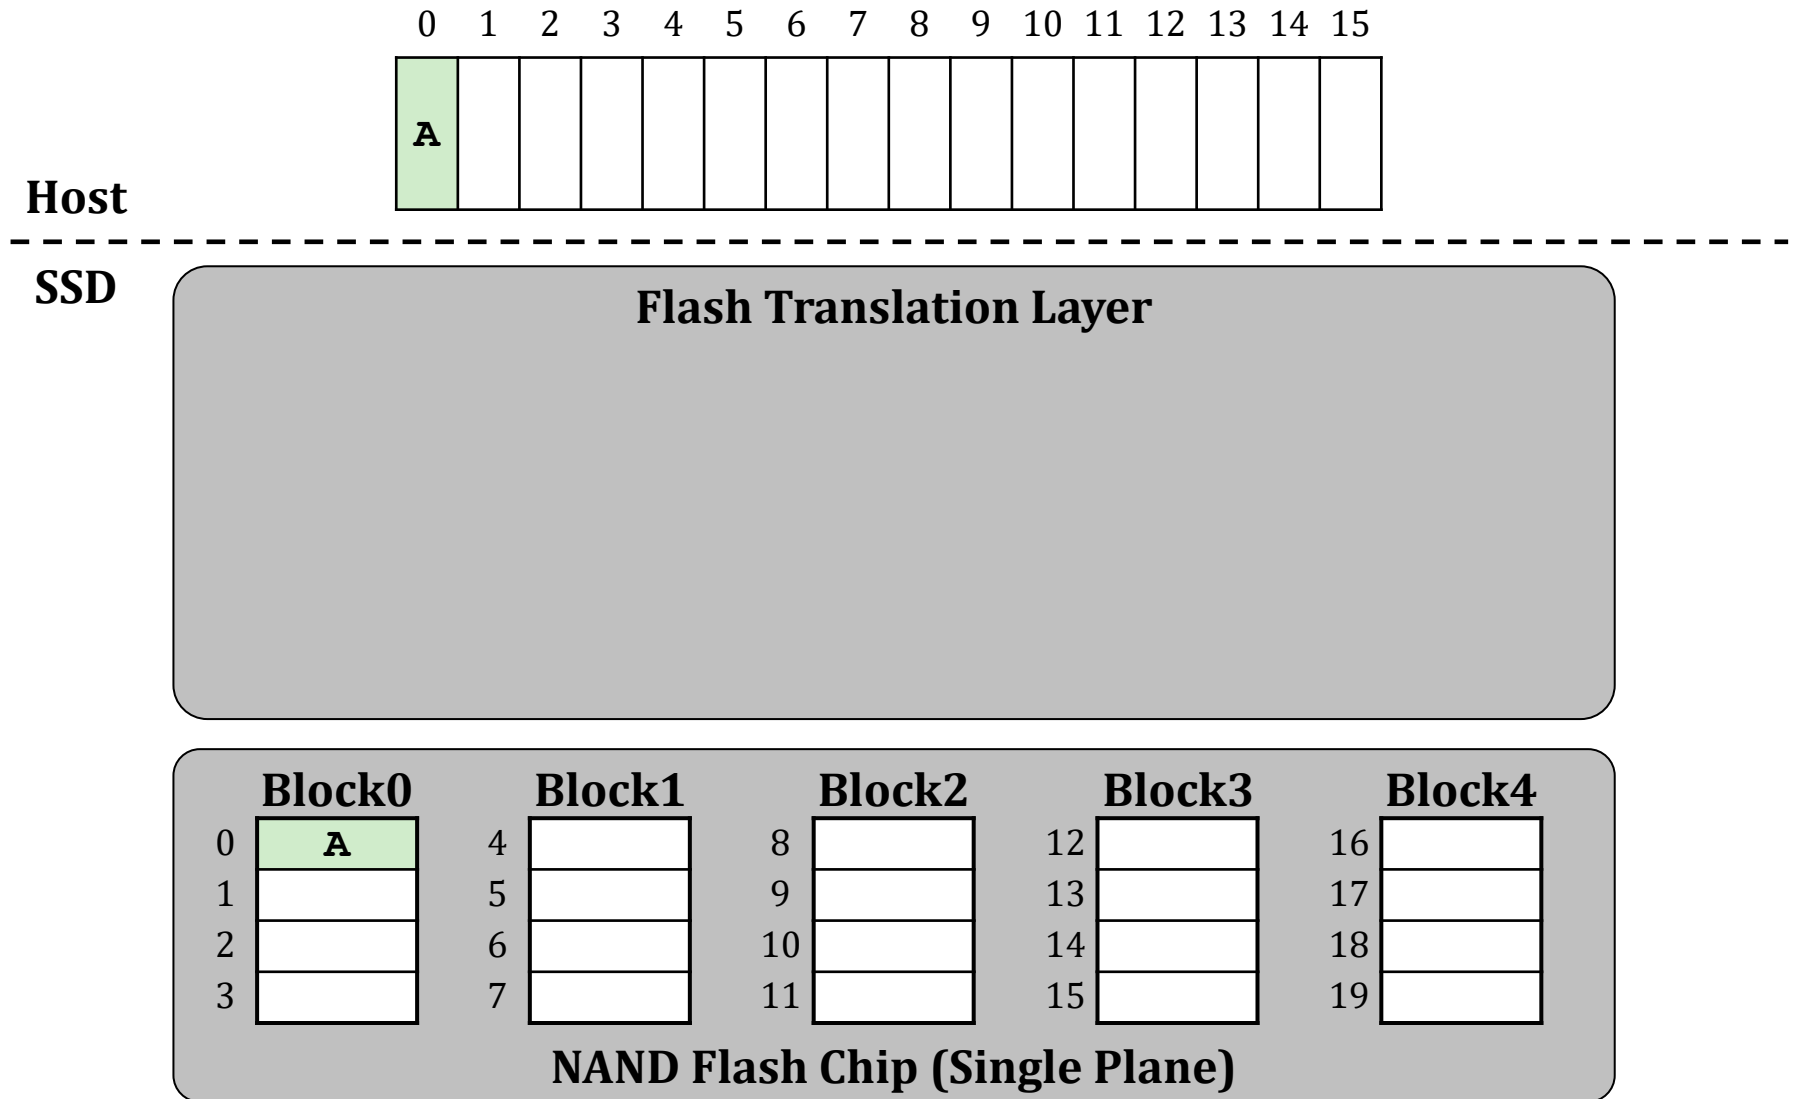
Write Request Handling: Page Write



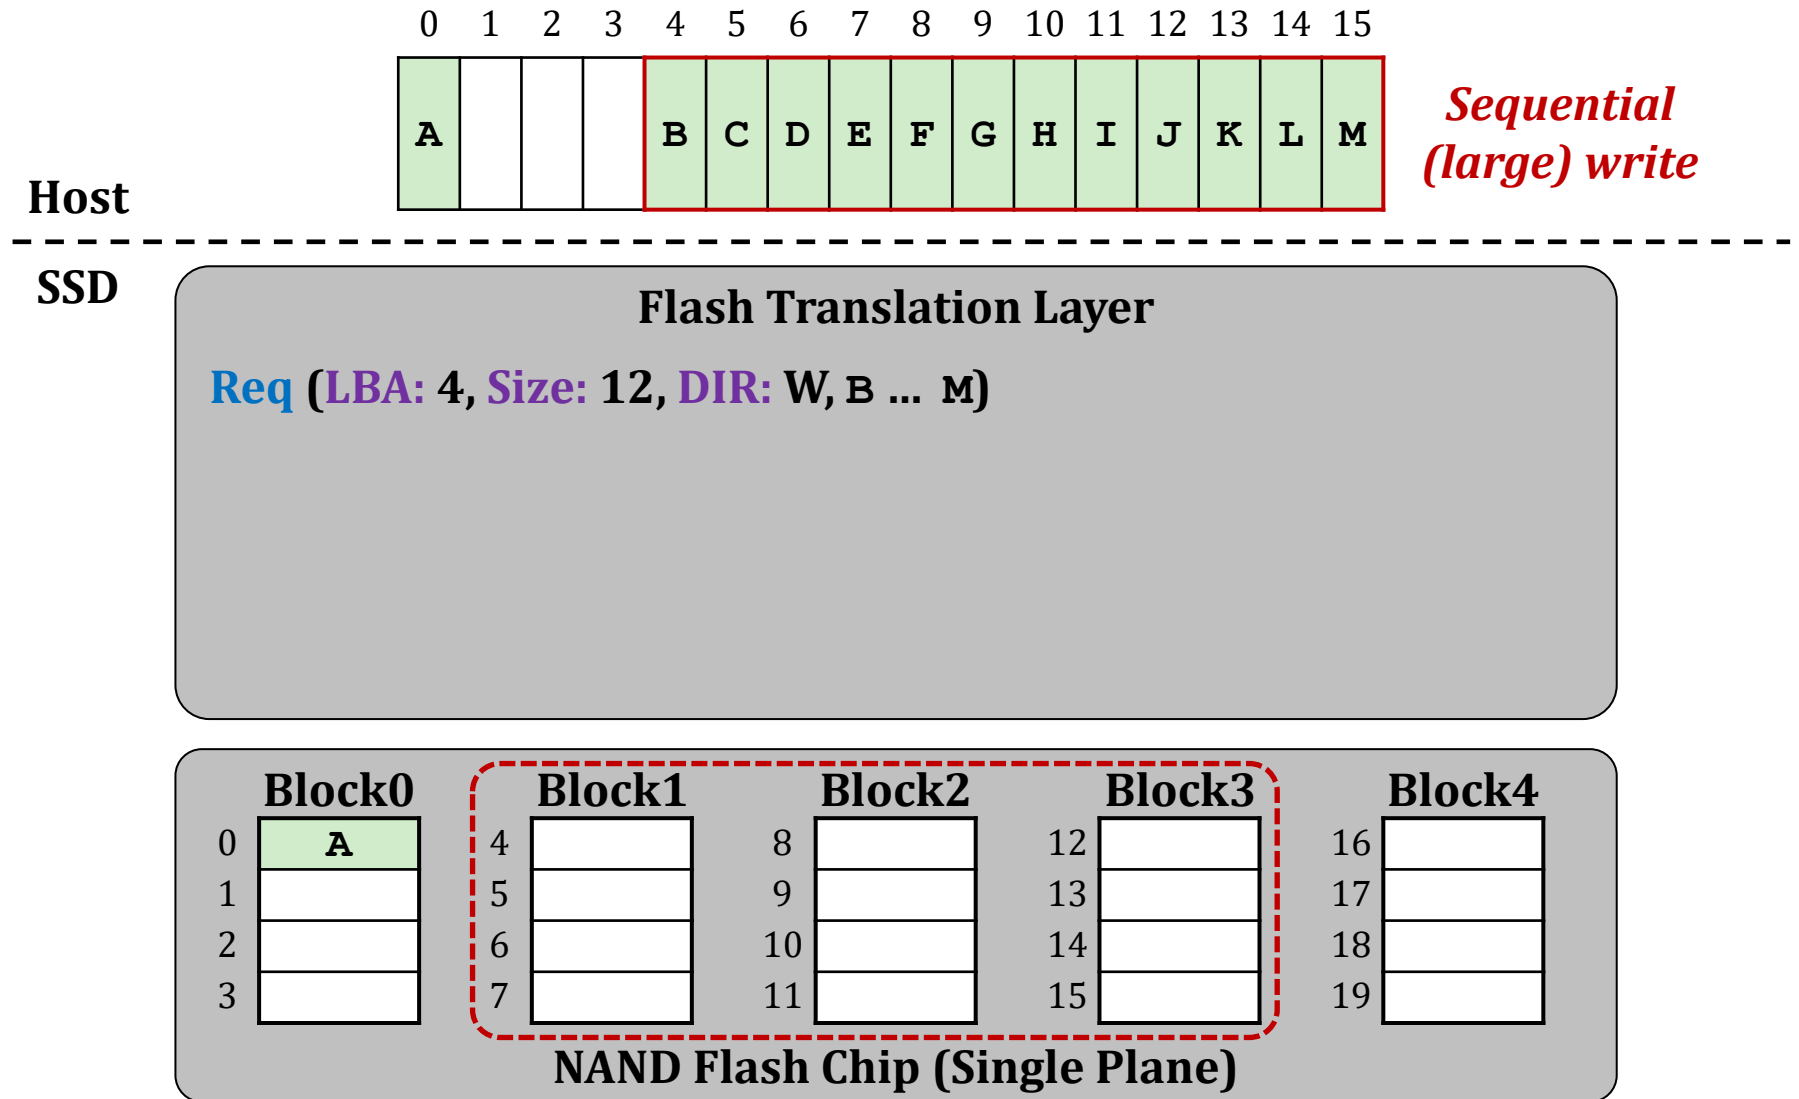
Write Request Handling: Page Write



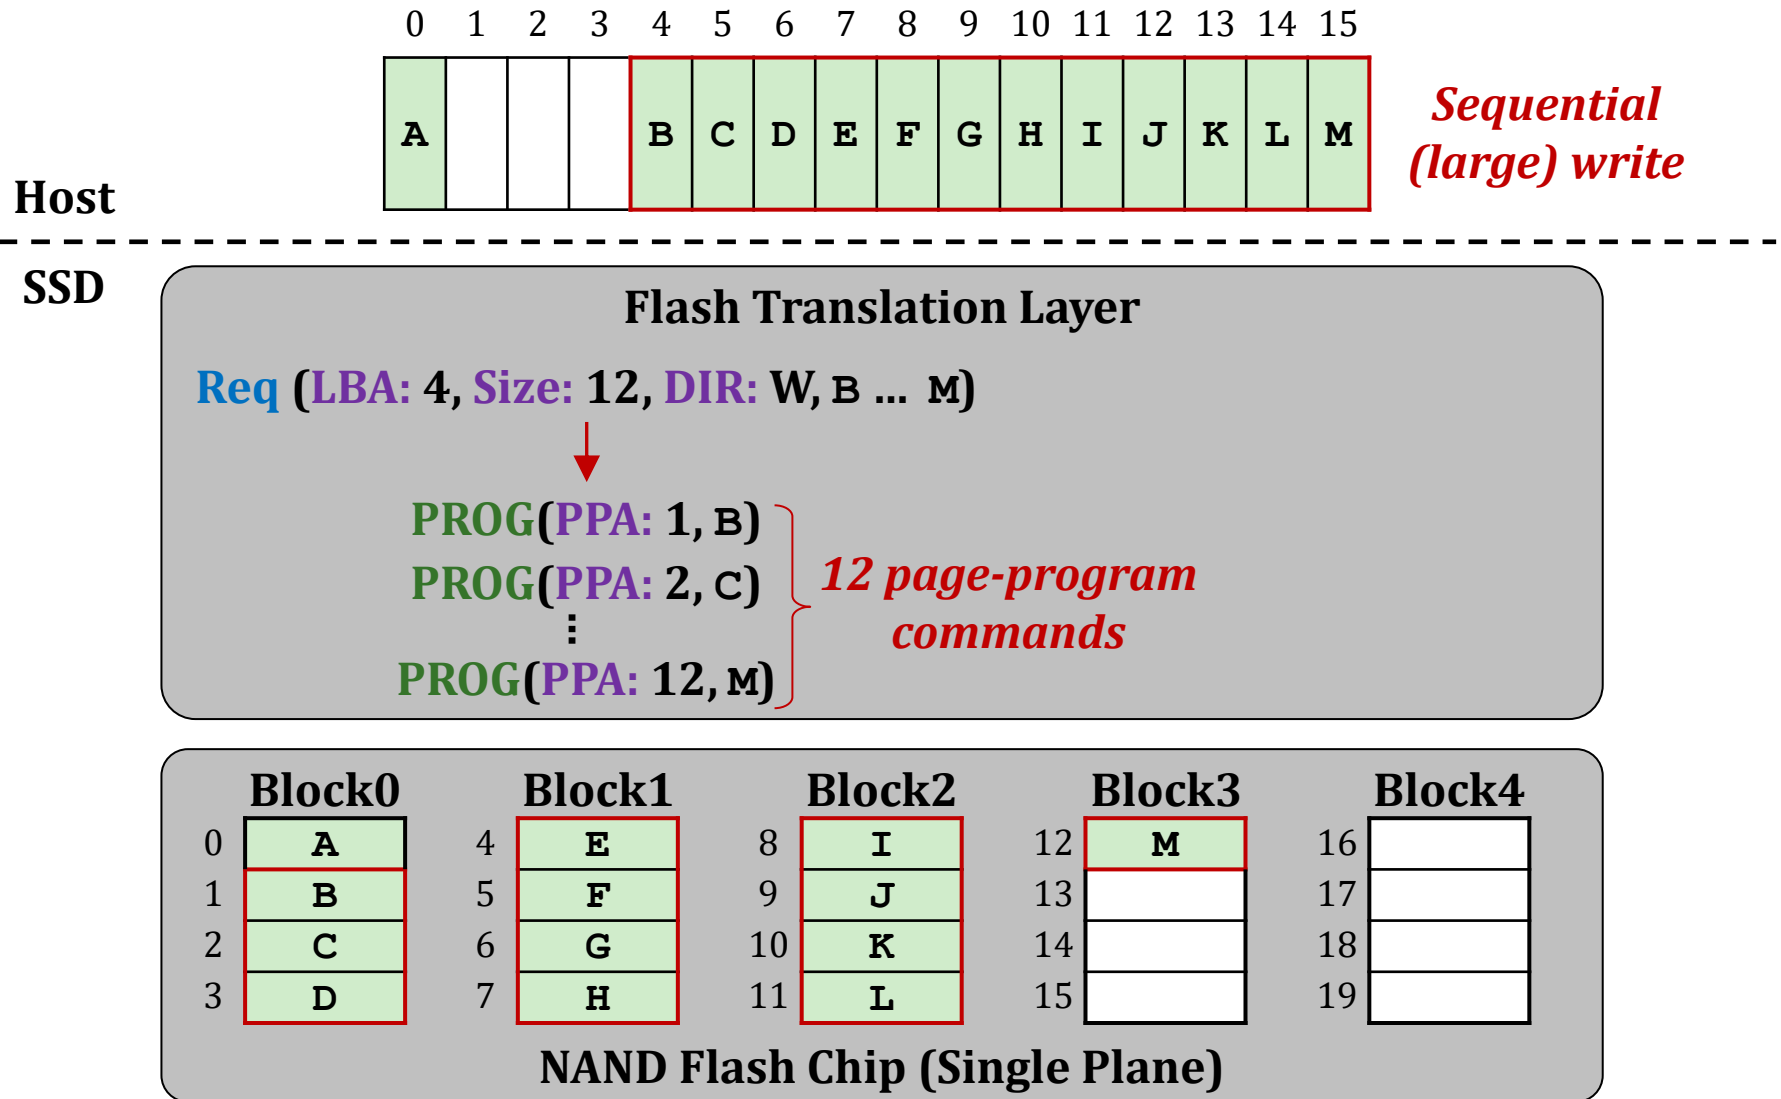
Write Request Handling: Sequential Write



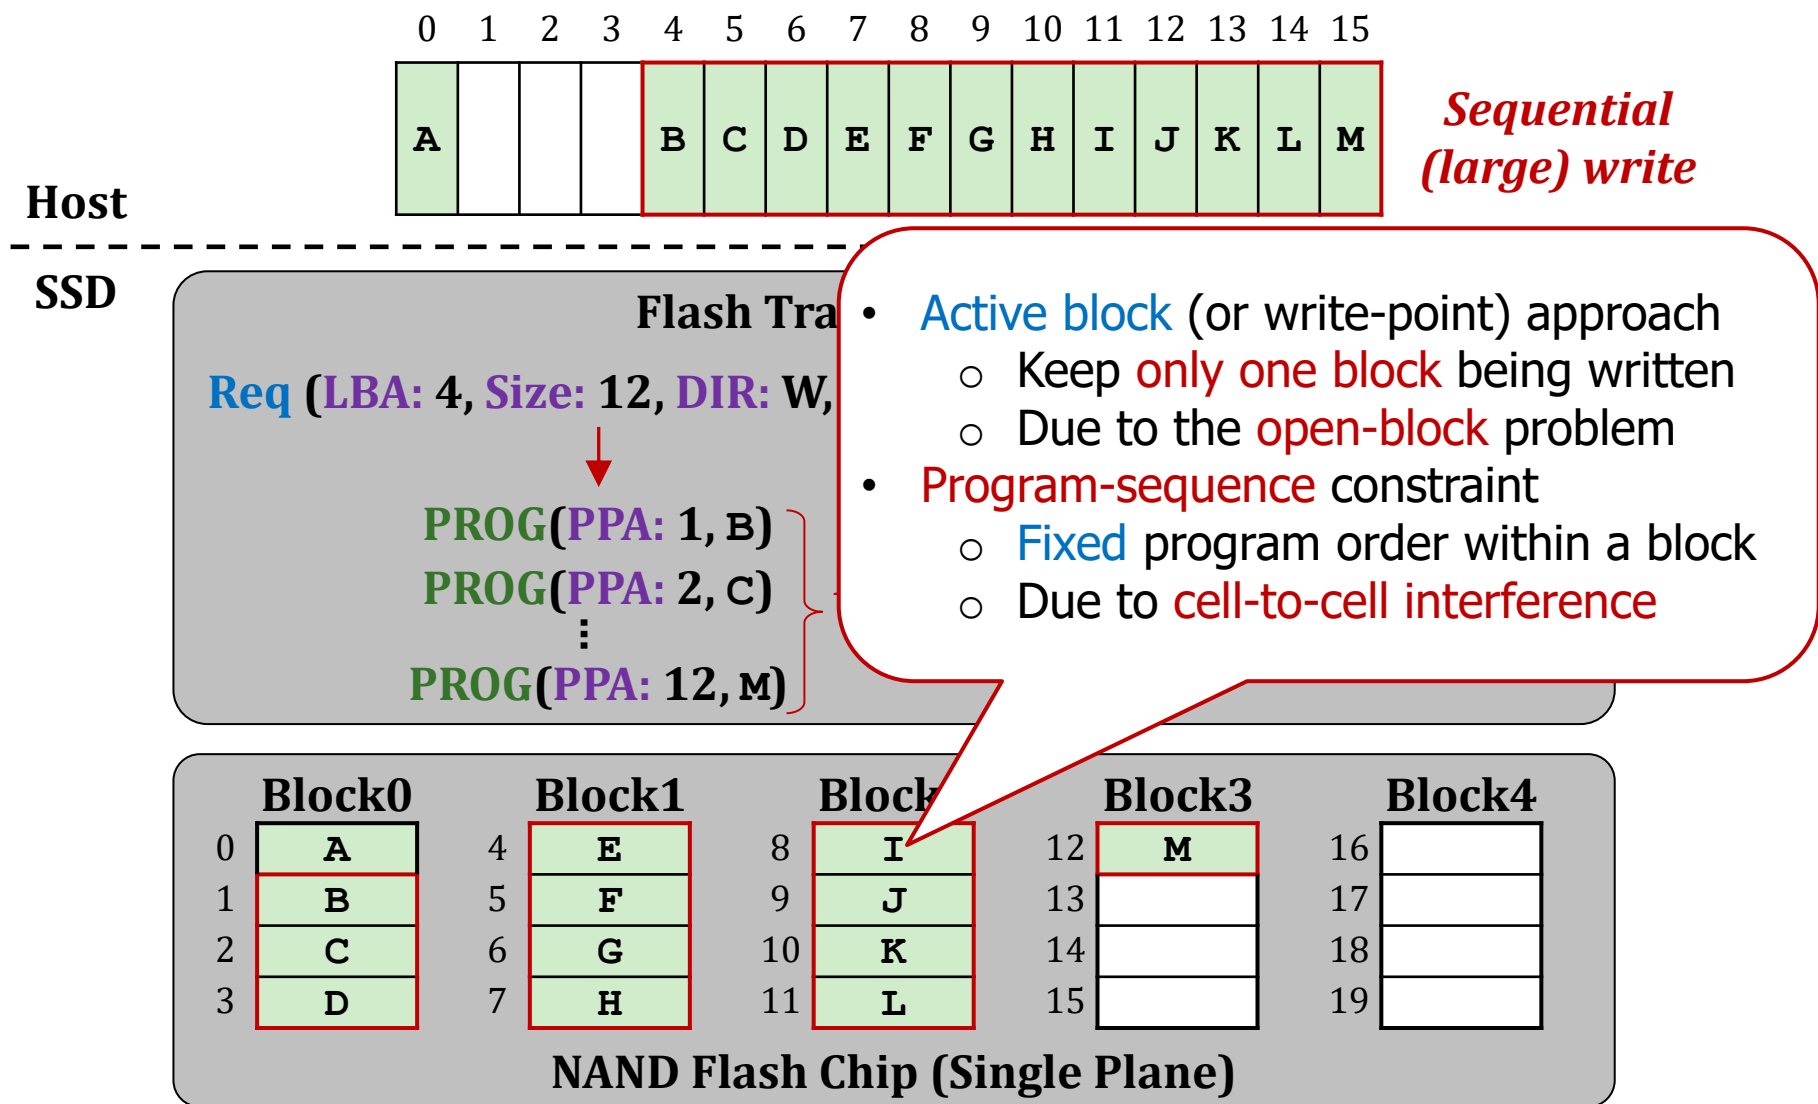
Write Request Handling: Sequential Write



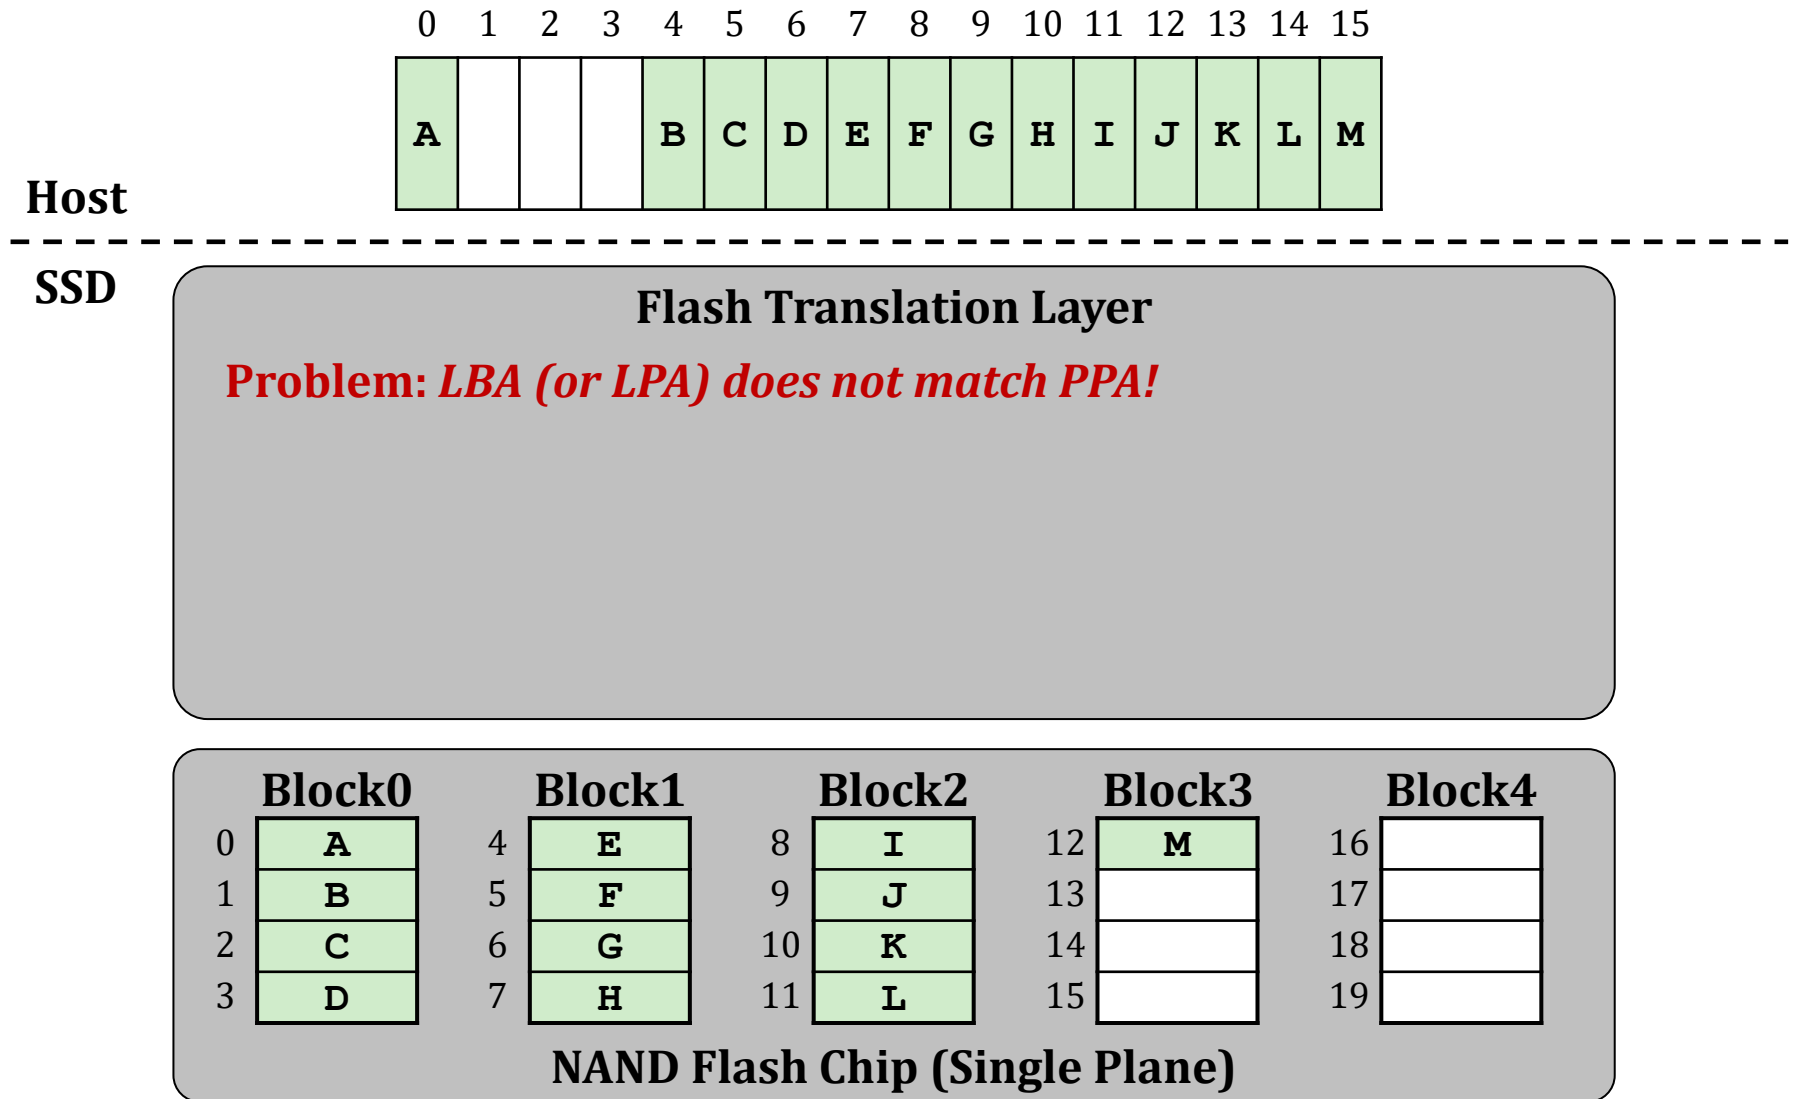
Write Request Handling: Sequential Write




Write Request Handling: Sequential Write



Write Request Handling: Address Mapping



Write Request Handling: Address Mapping



SSD


Flash Translation Layer

Problem: LBA (or LPA) does not match PPA!

Req (LBA: 4, Size: 1, DIR: R)

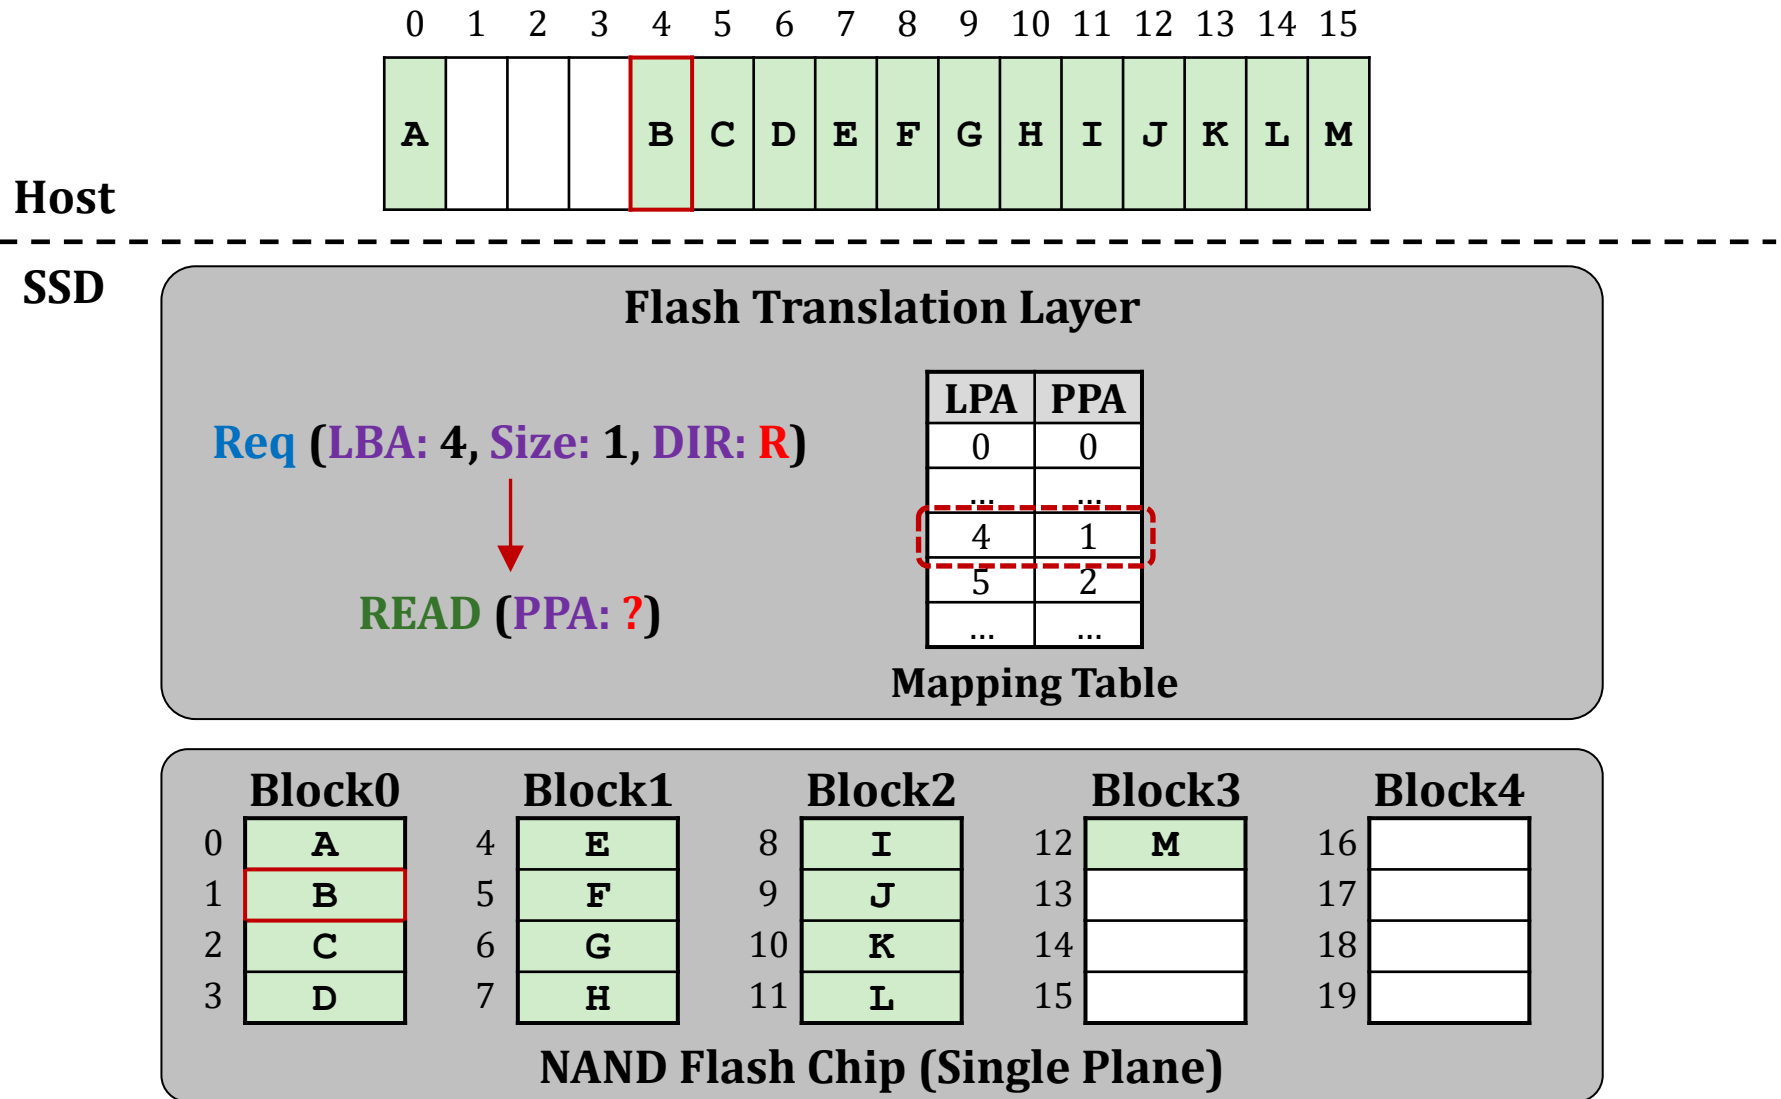
↓
READ (PPA: ?)

*Needs to maintain
Address-mapping information*

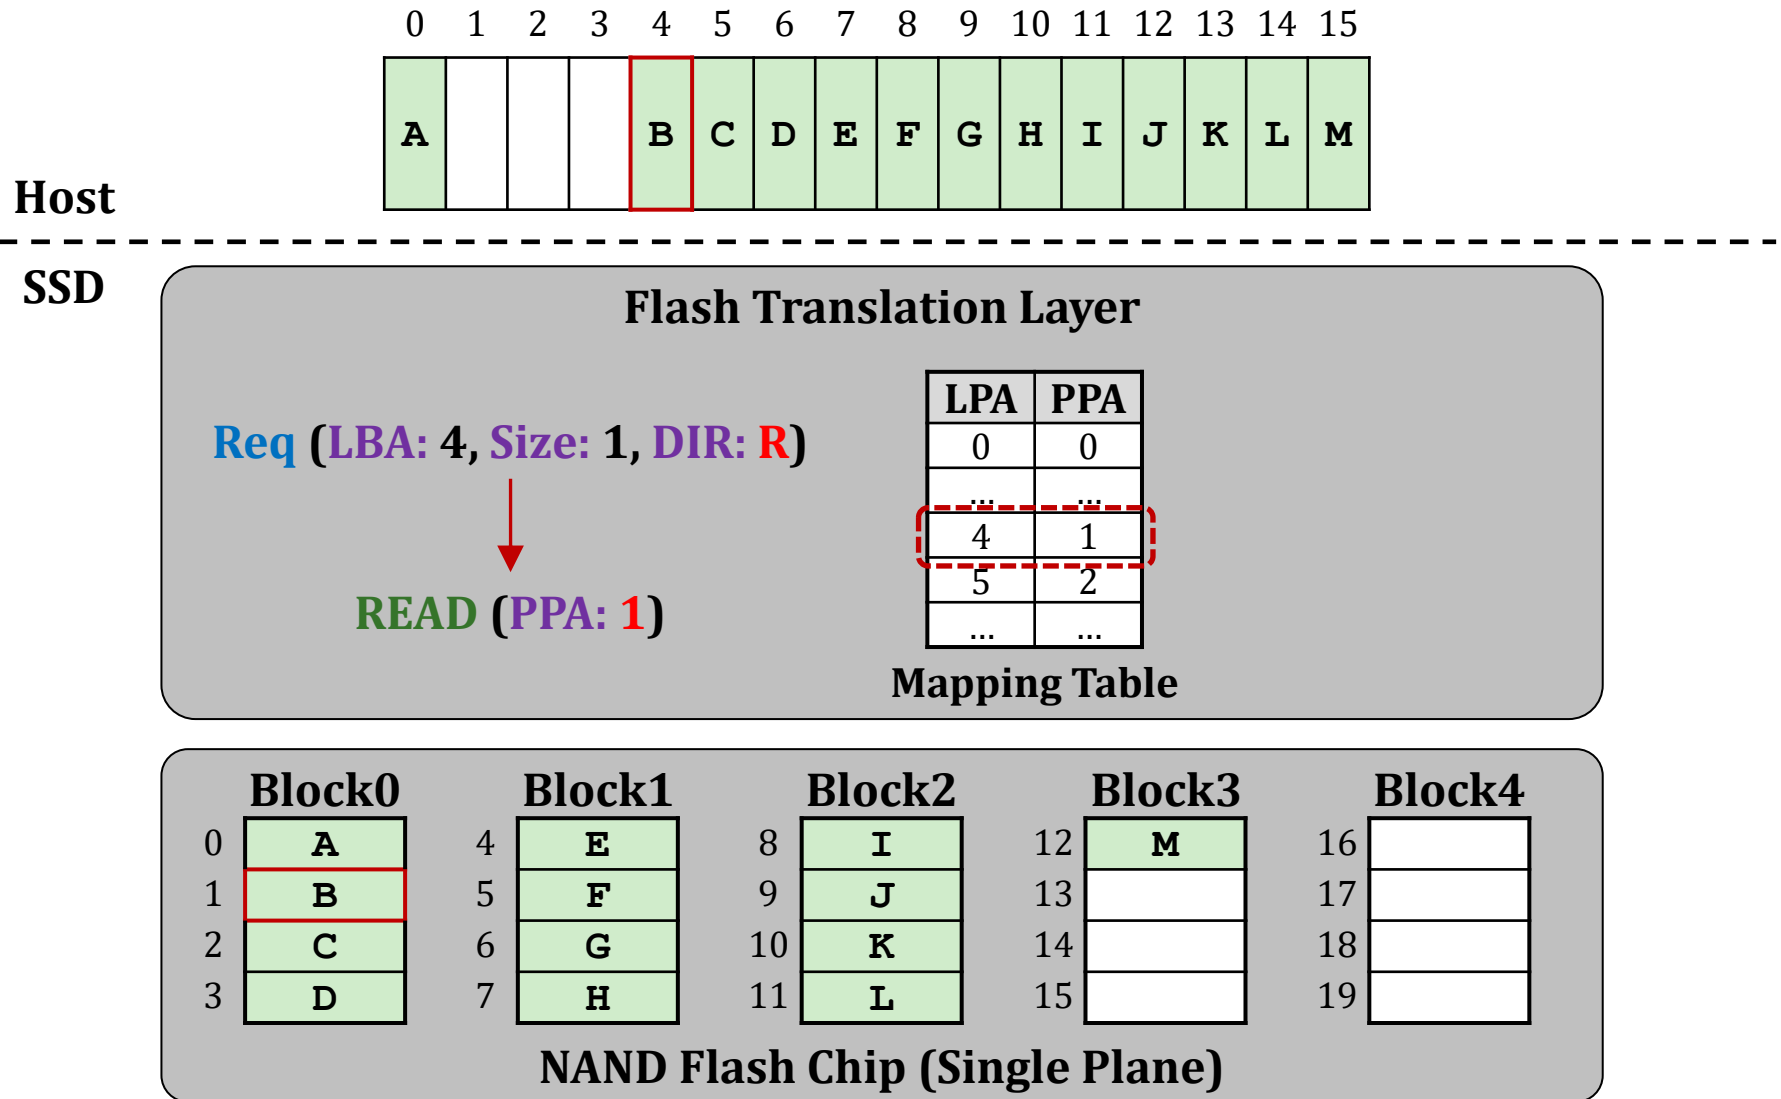


NAND Flash Chip (Single Plane)

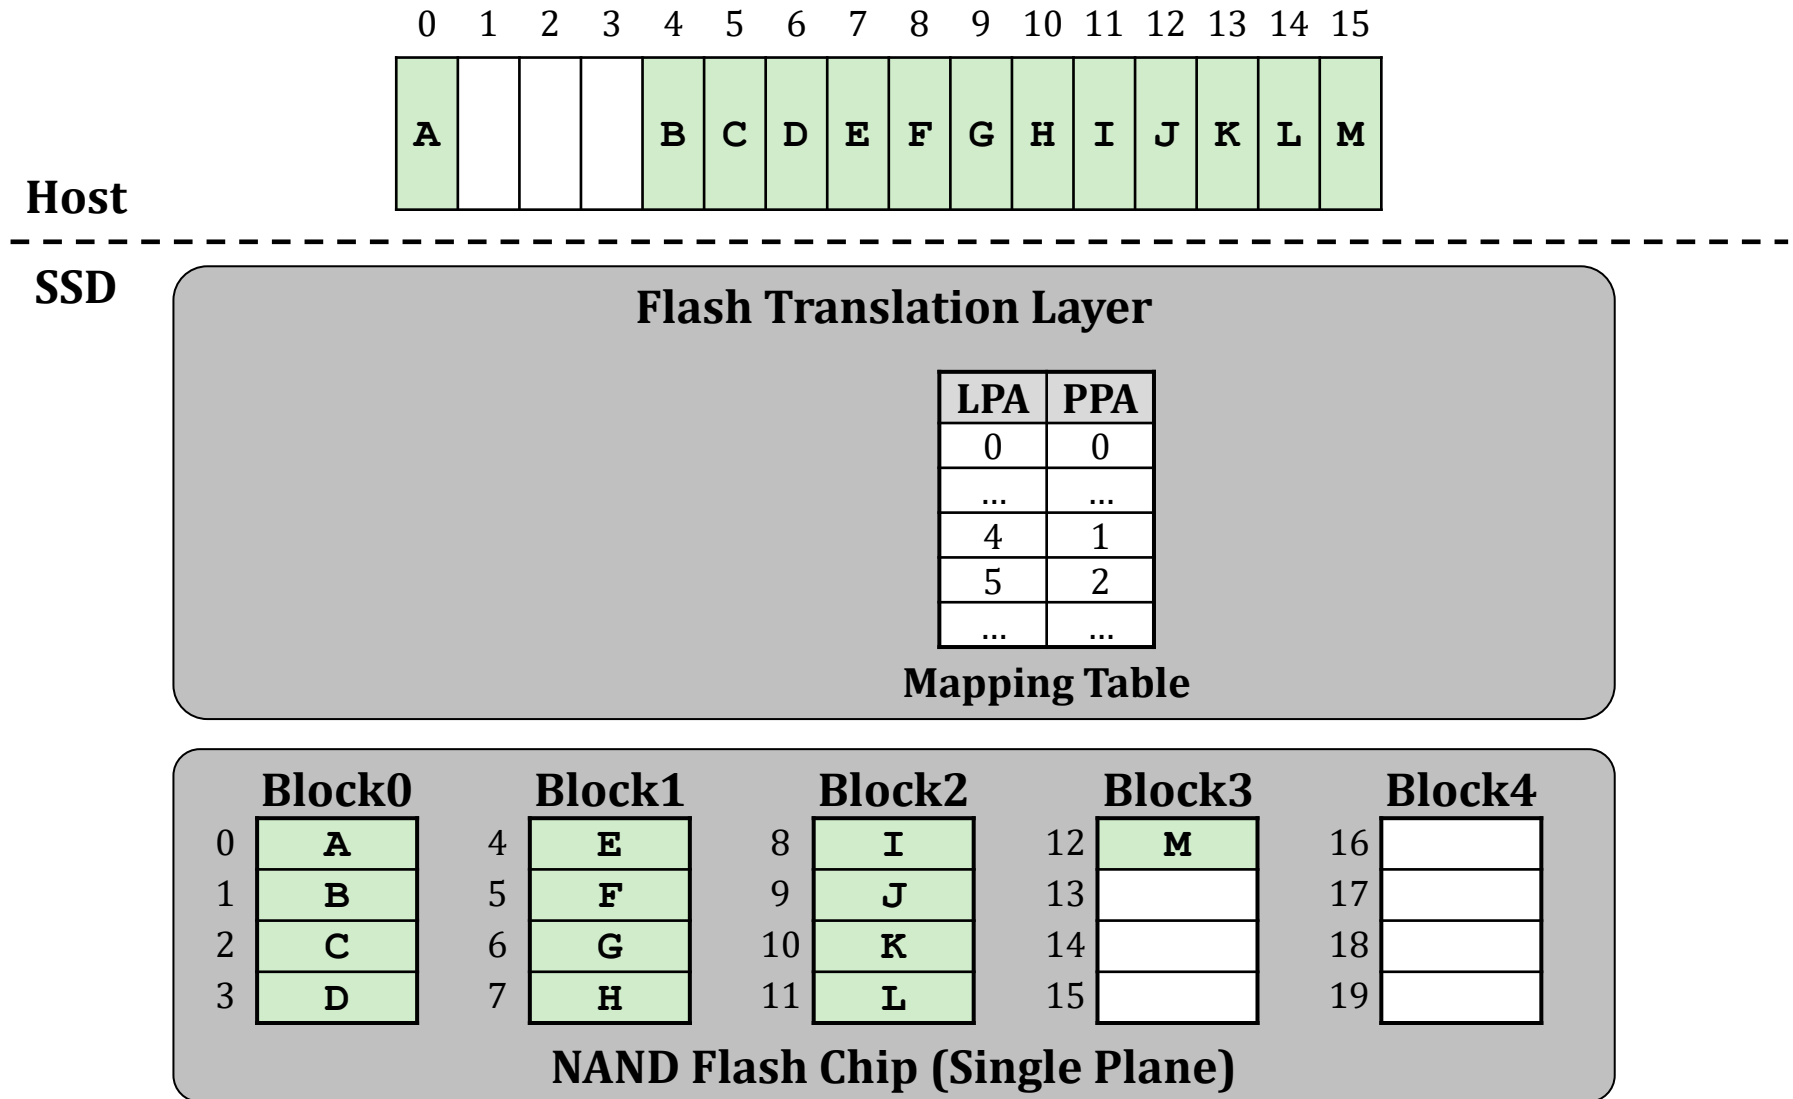
Write Request Handling: Address Mapping



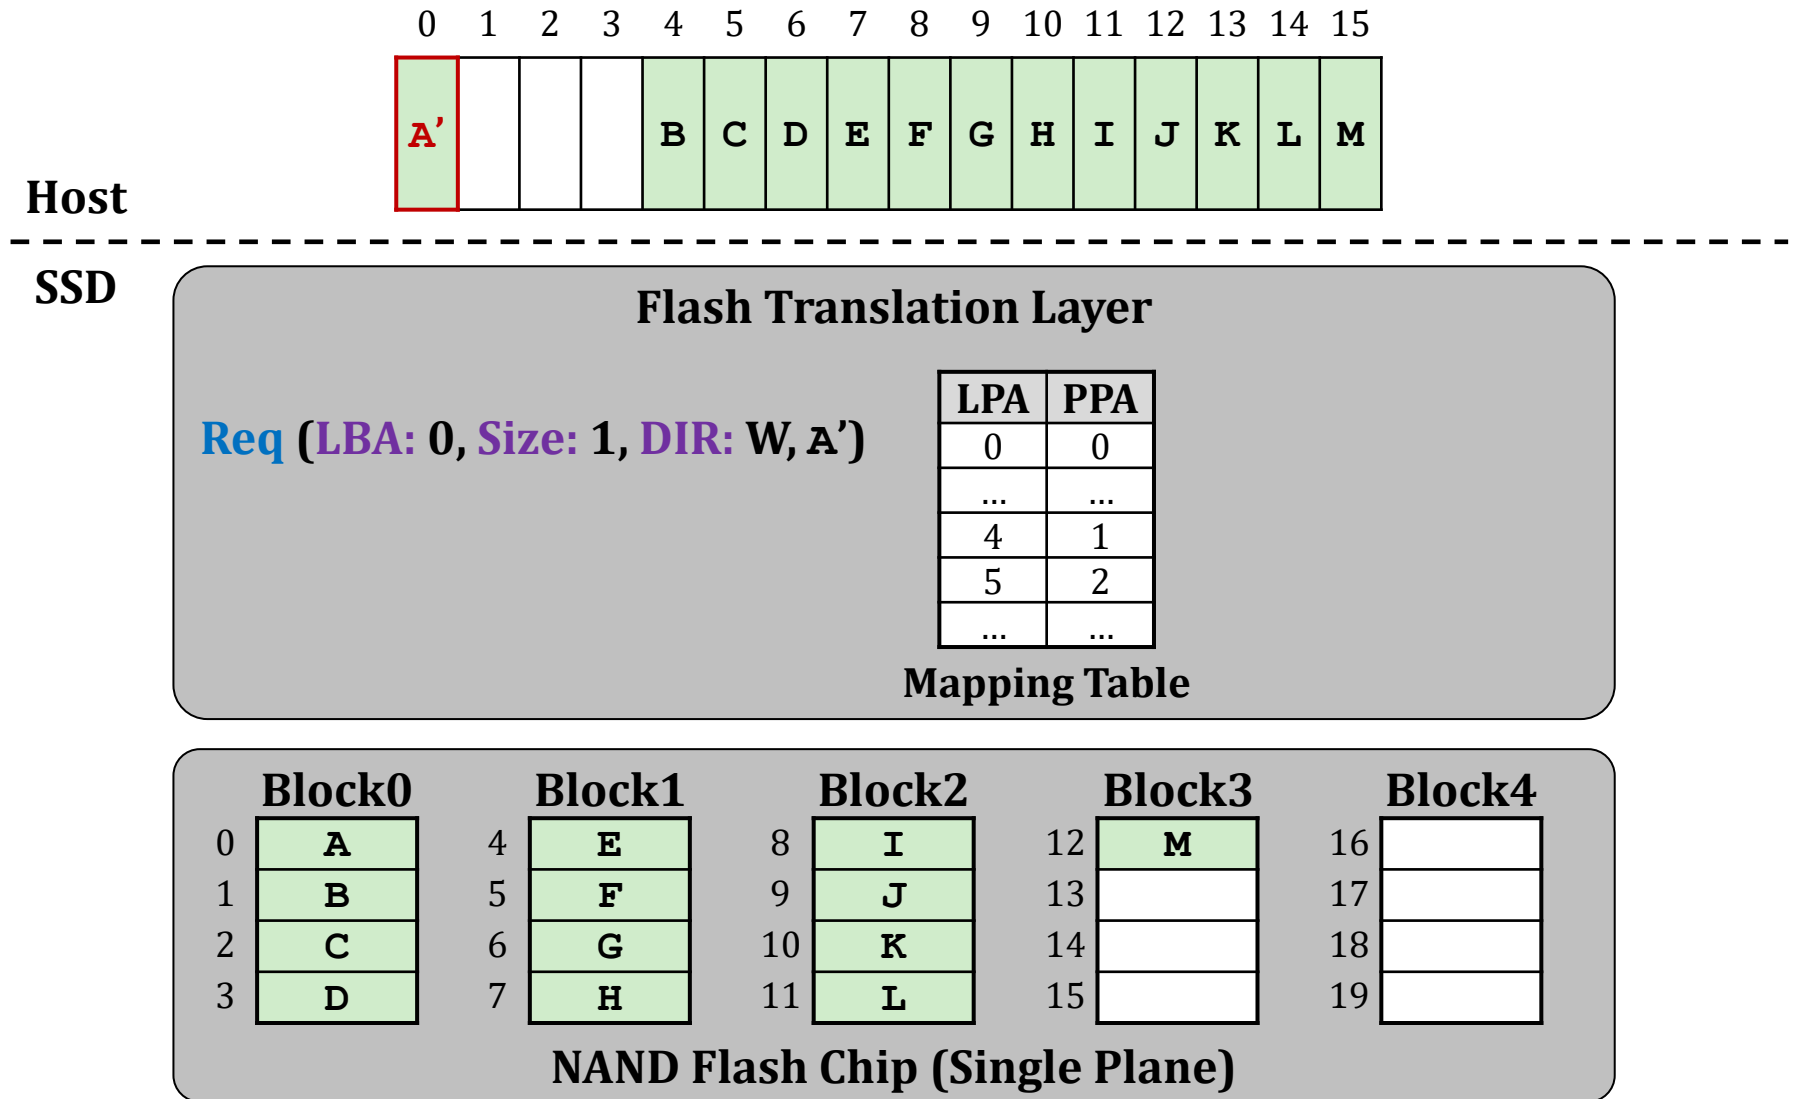
Write Request Handling: Address Mapping



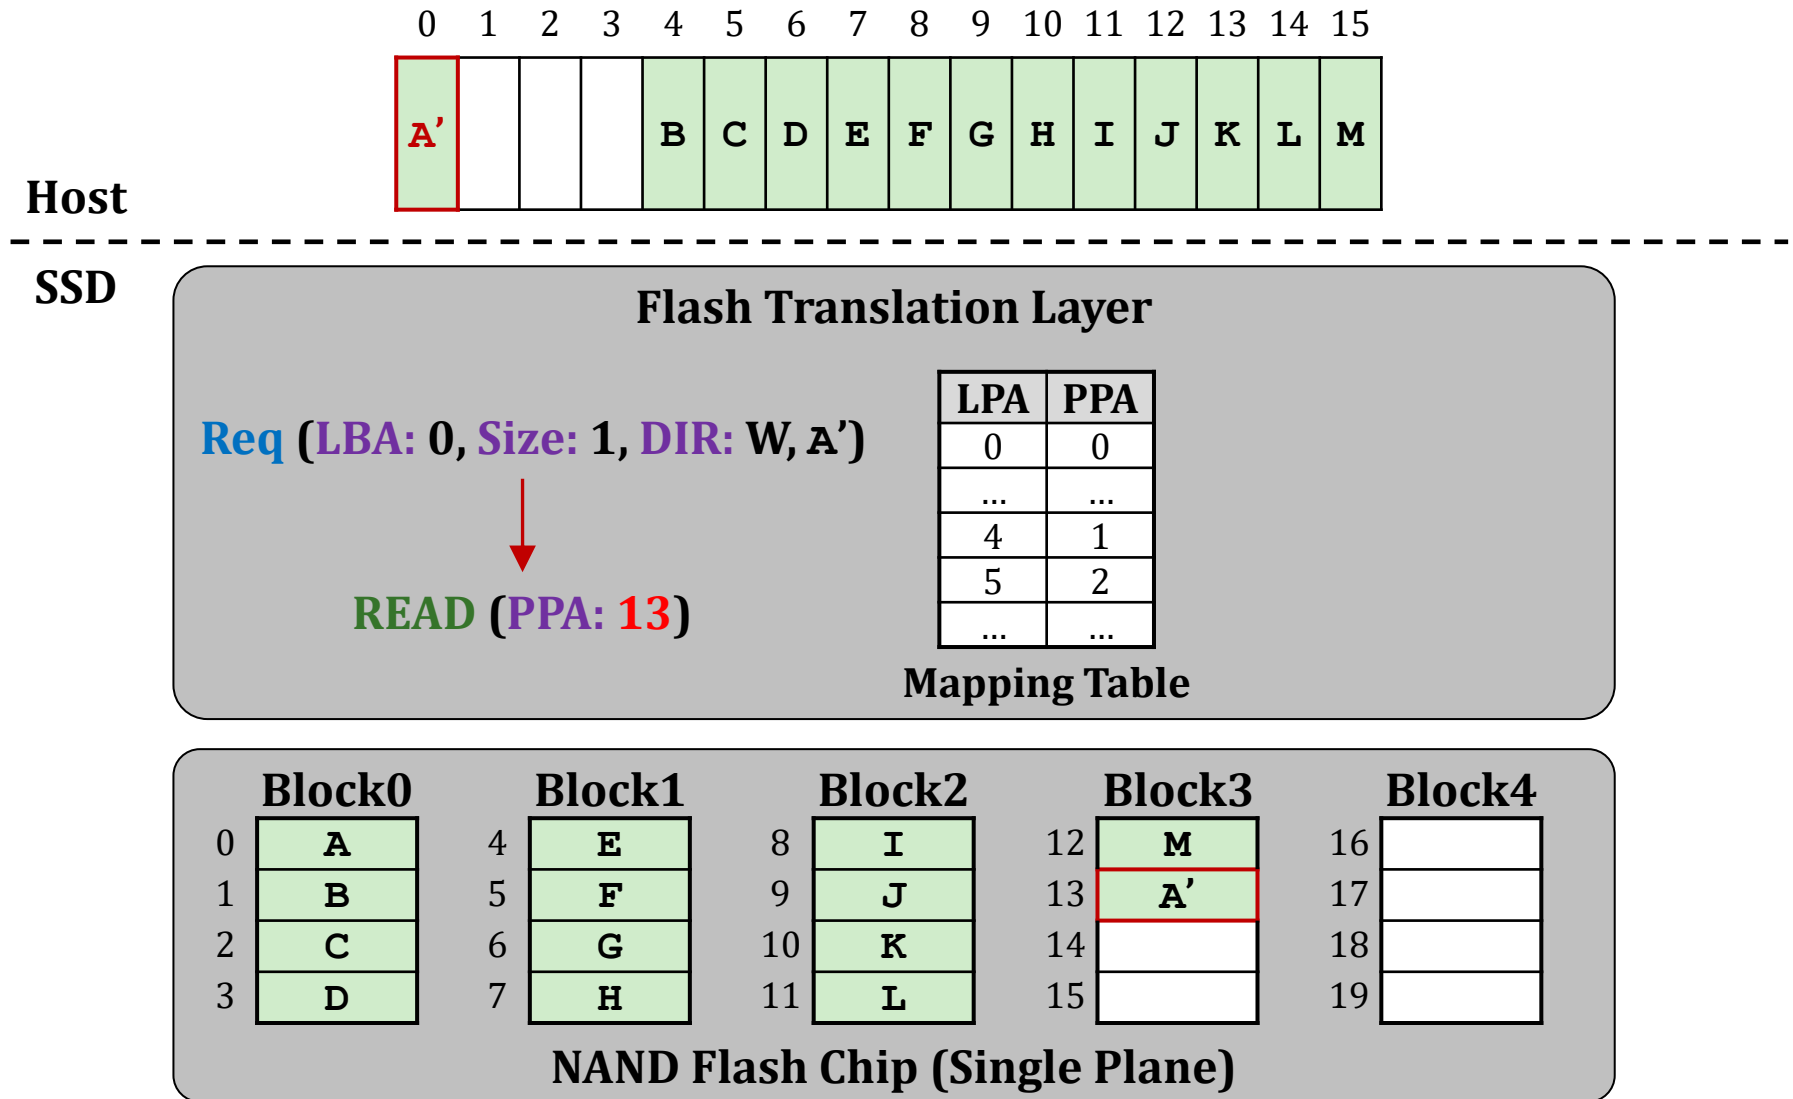
Write Request Handling: Update



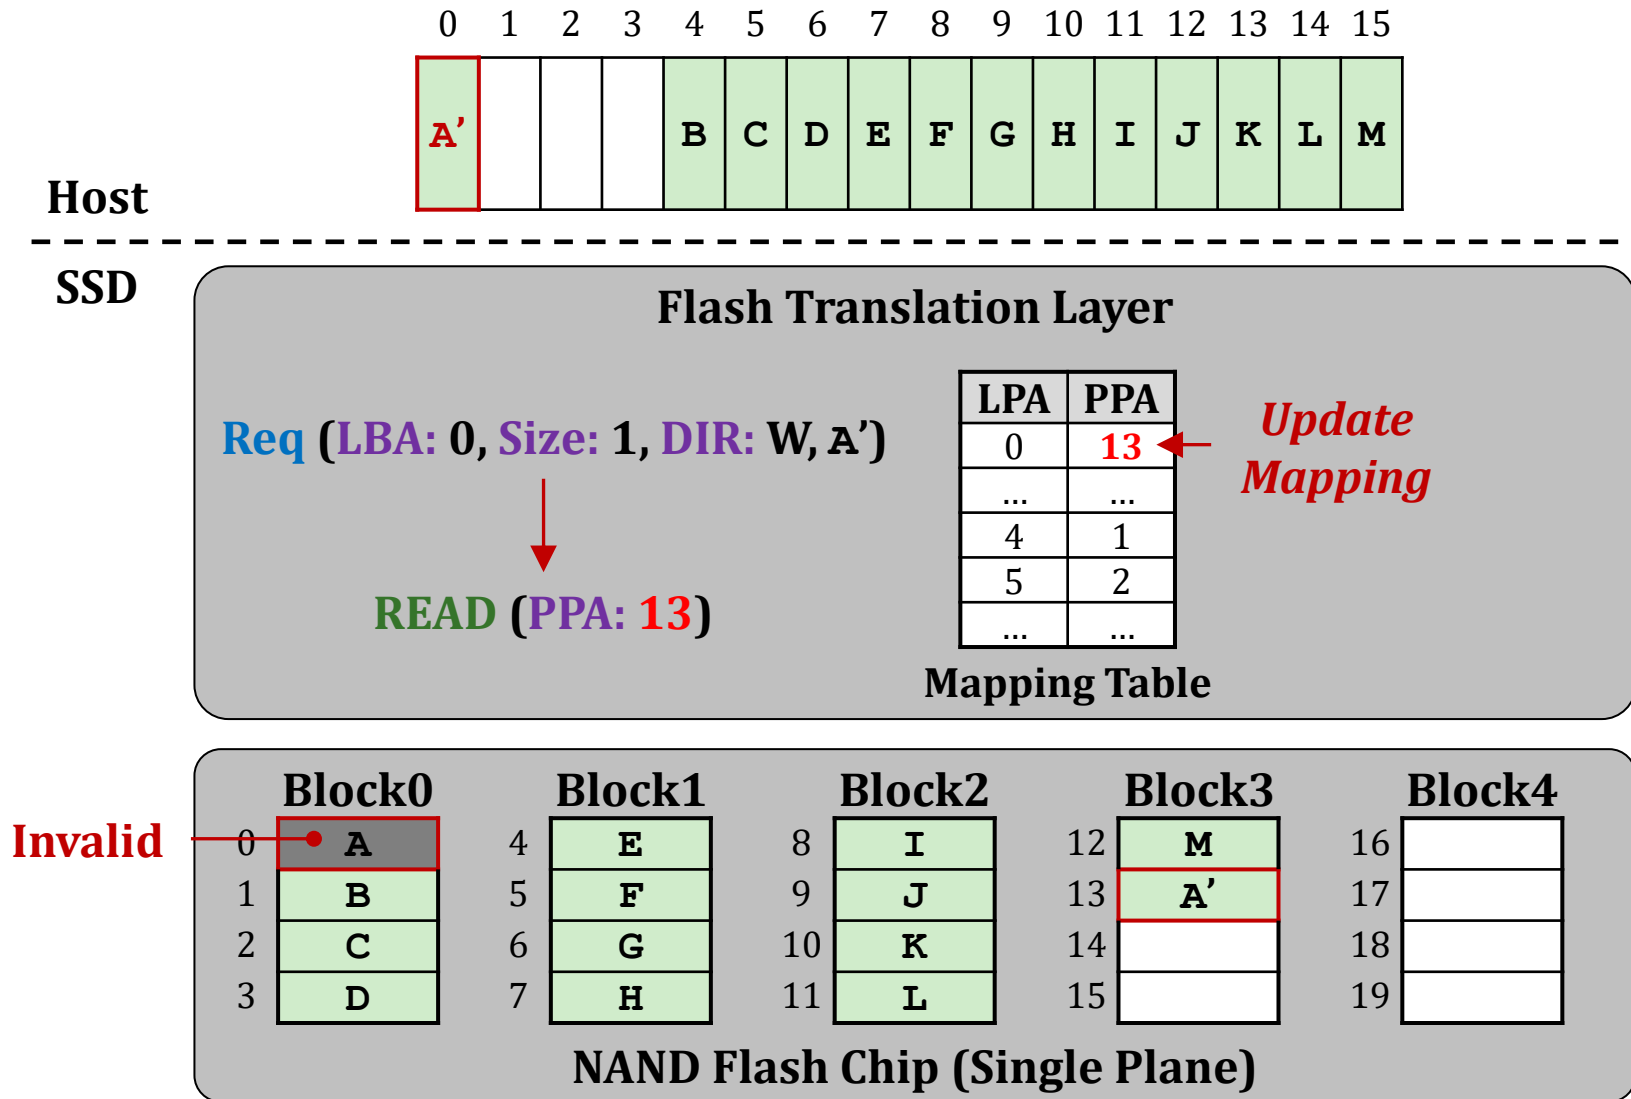
Write Request Handling: Update



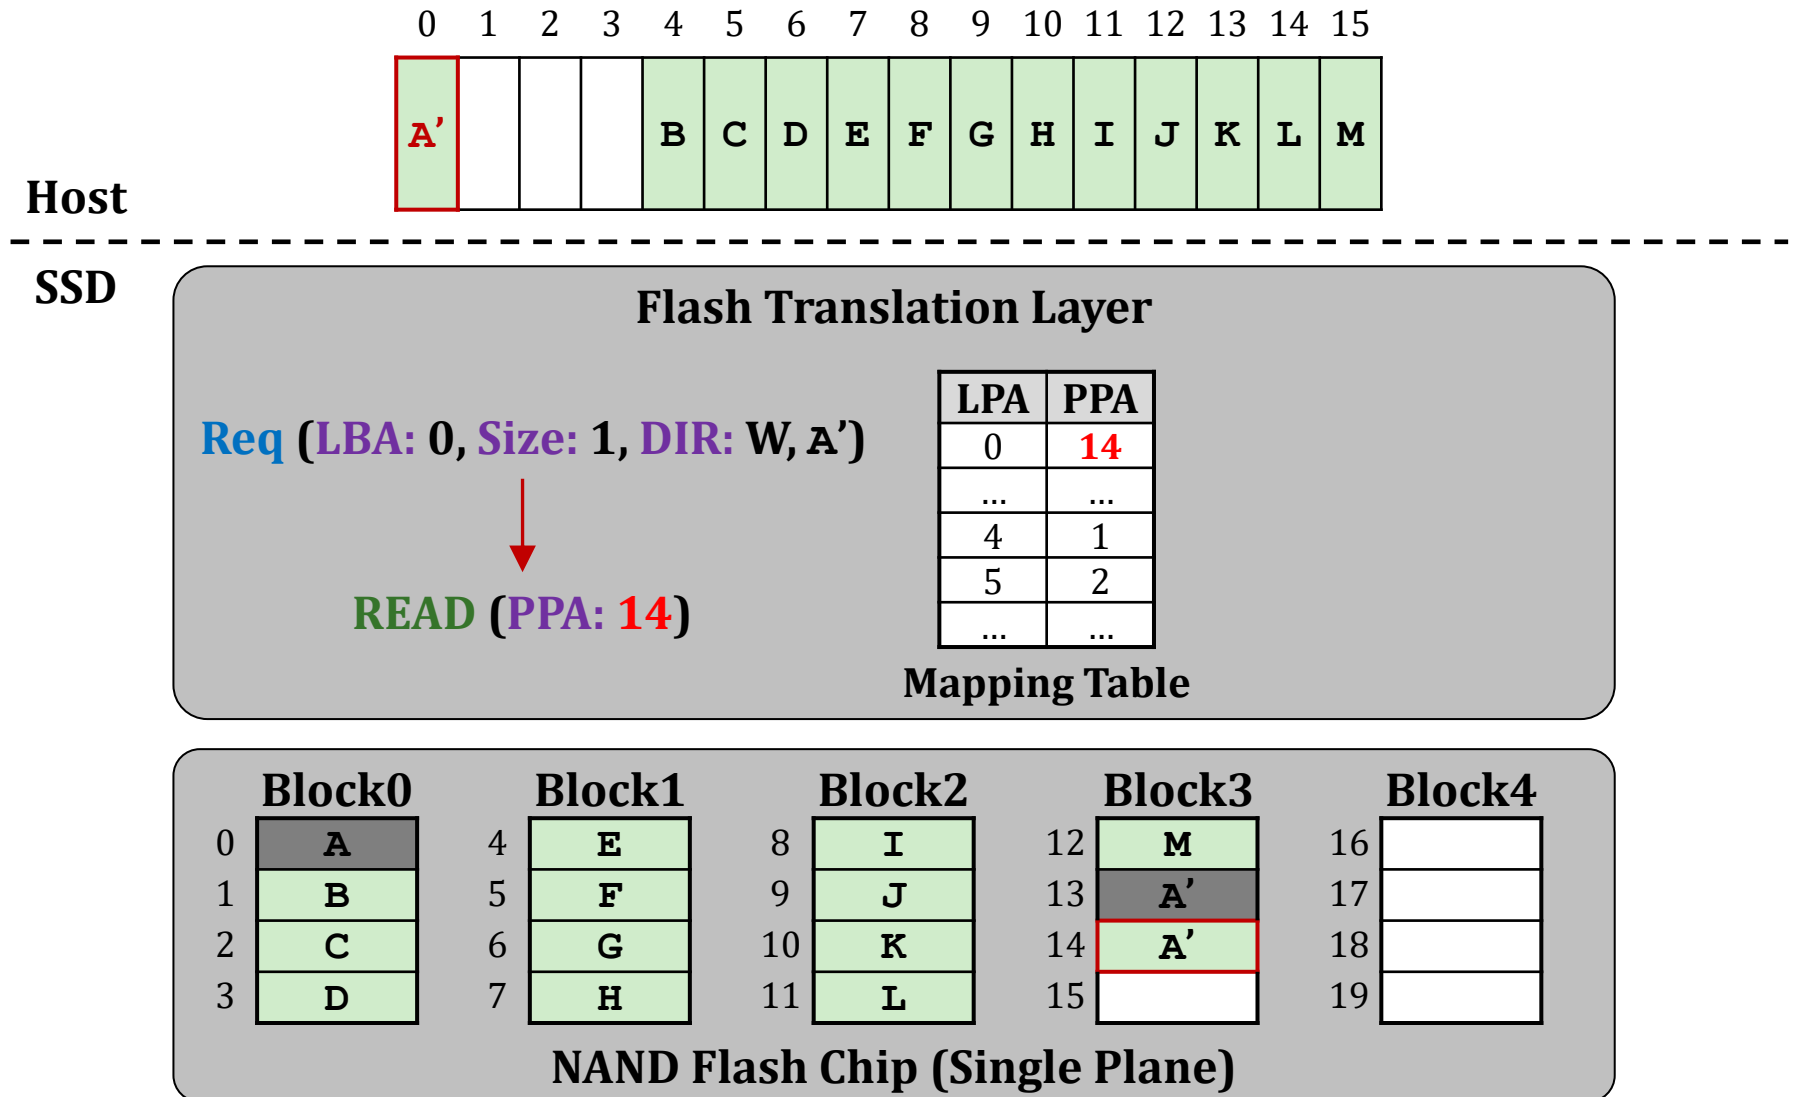
Write Request Handling: Update



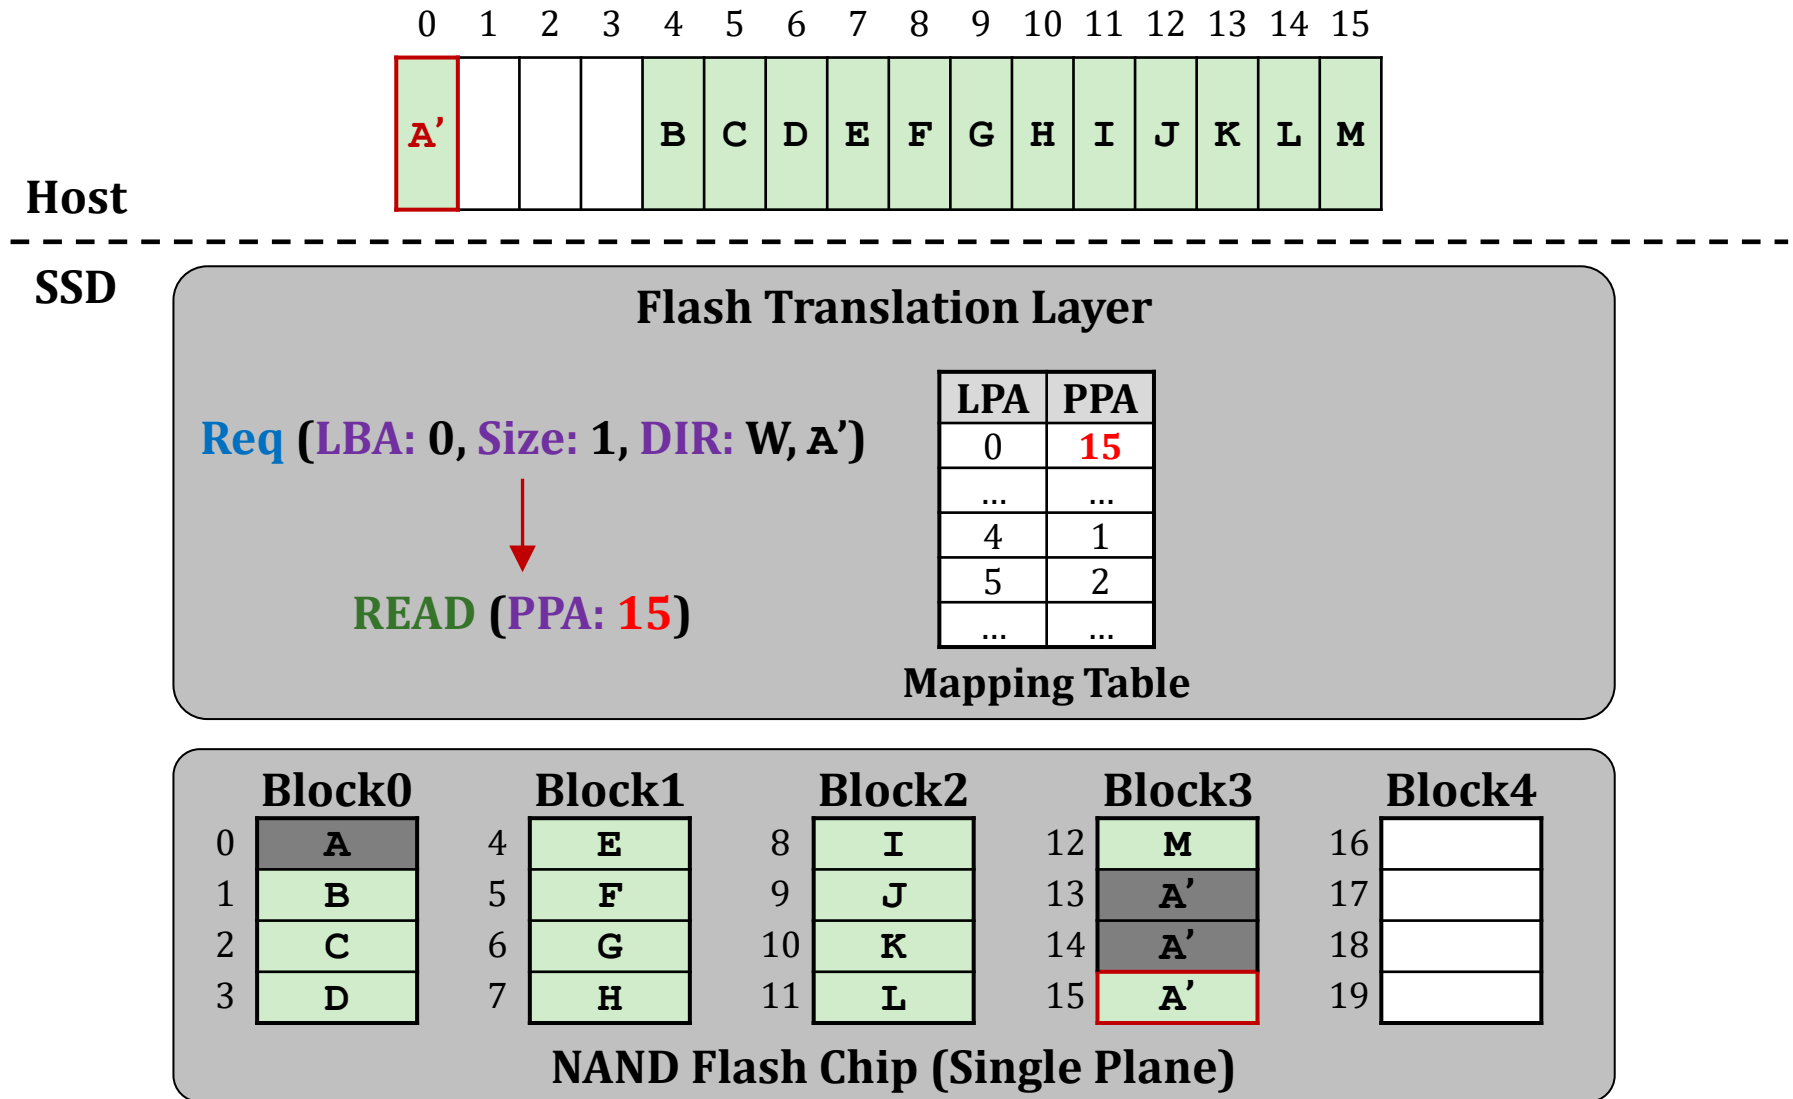
Write Request Handling: Update



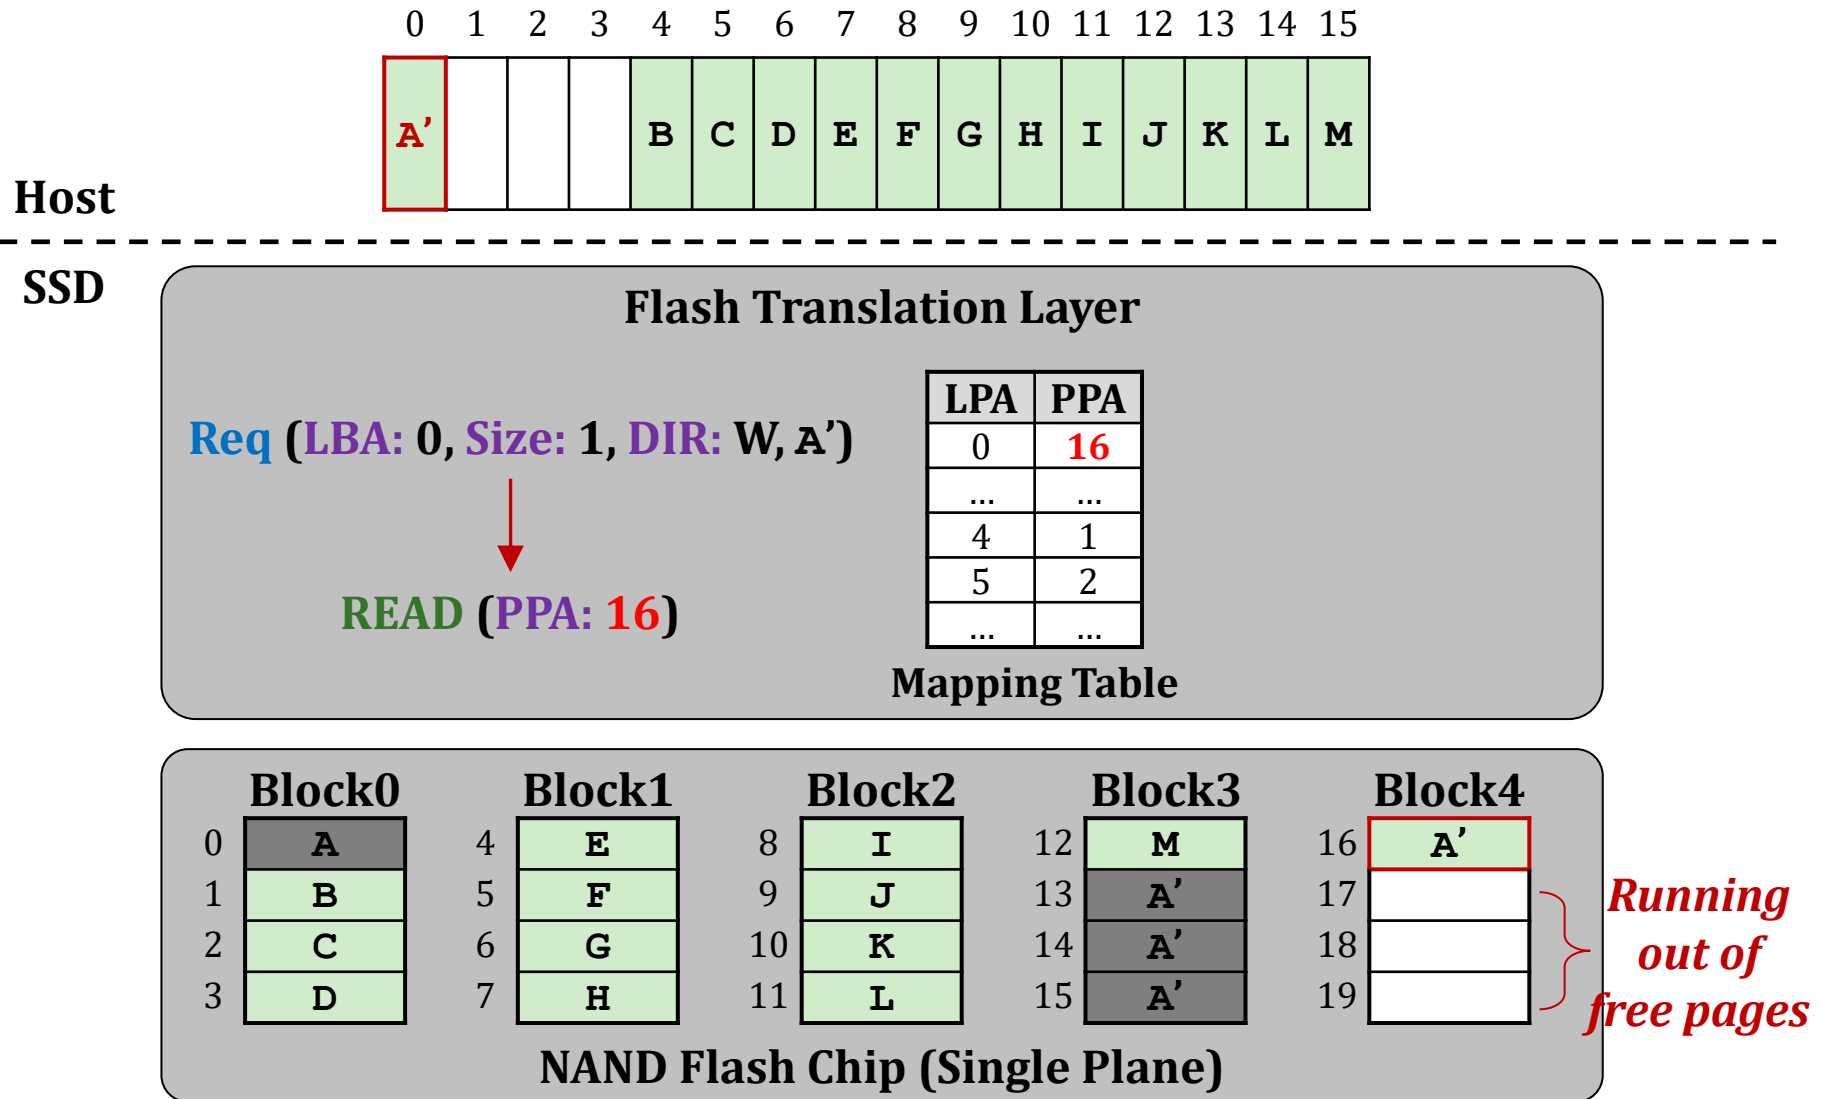
Write Request Handling: Update



Write Request Handling: Update



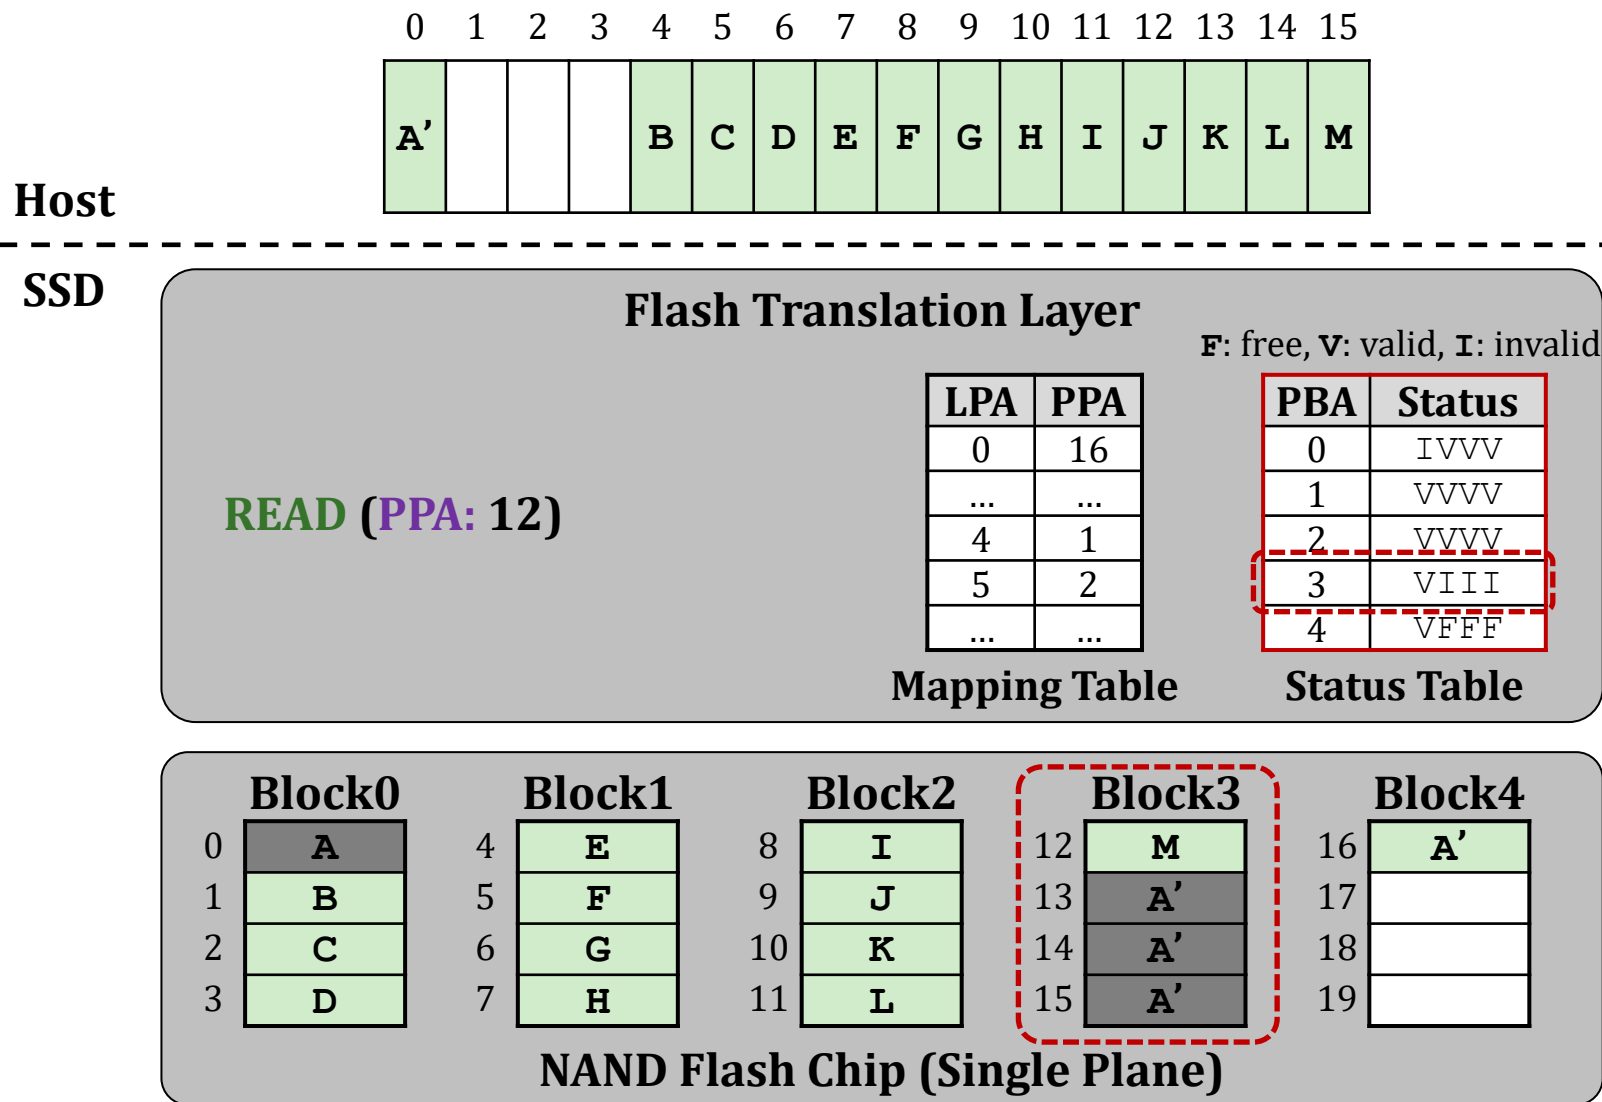
Write Request Handling: Update



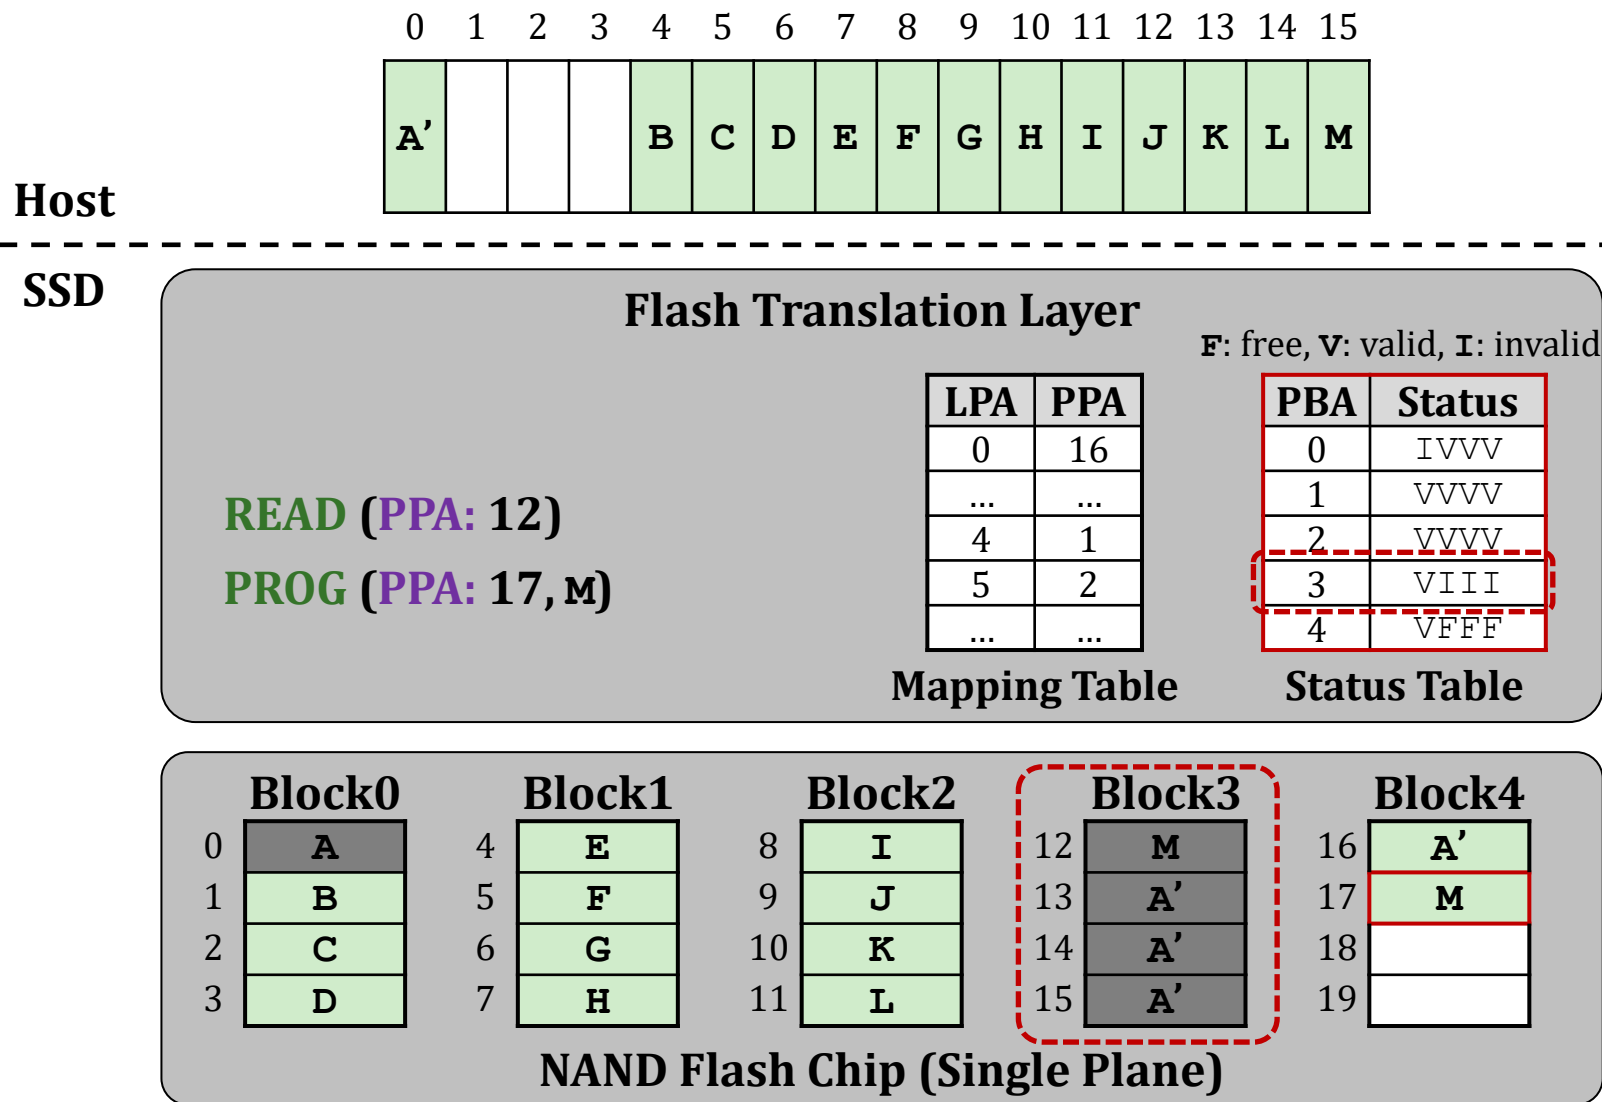
Garbage Collection

- Reclaims **free pages** by erasing **invalid** pages
 - Erase unit: **block**
 - If a victim block (to erase) has **valid pages**,
all the valid pages **need to be copied** to other free pages
 - **Performance overhead:** $(t_{\text{READ}} + t_{\text{PROG}}) \times \# \text{ of valid pages}$
 - **Lifetime overhead:** additional writes \rightarrow P/E-cycle increase
- **Greedy** victim-selection policy:
Erases the block with the **largest number** of invalid pages
 - Needs to maintain **# of invalid (or valid) pages** for each block

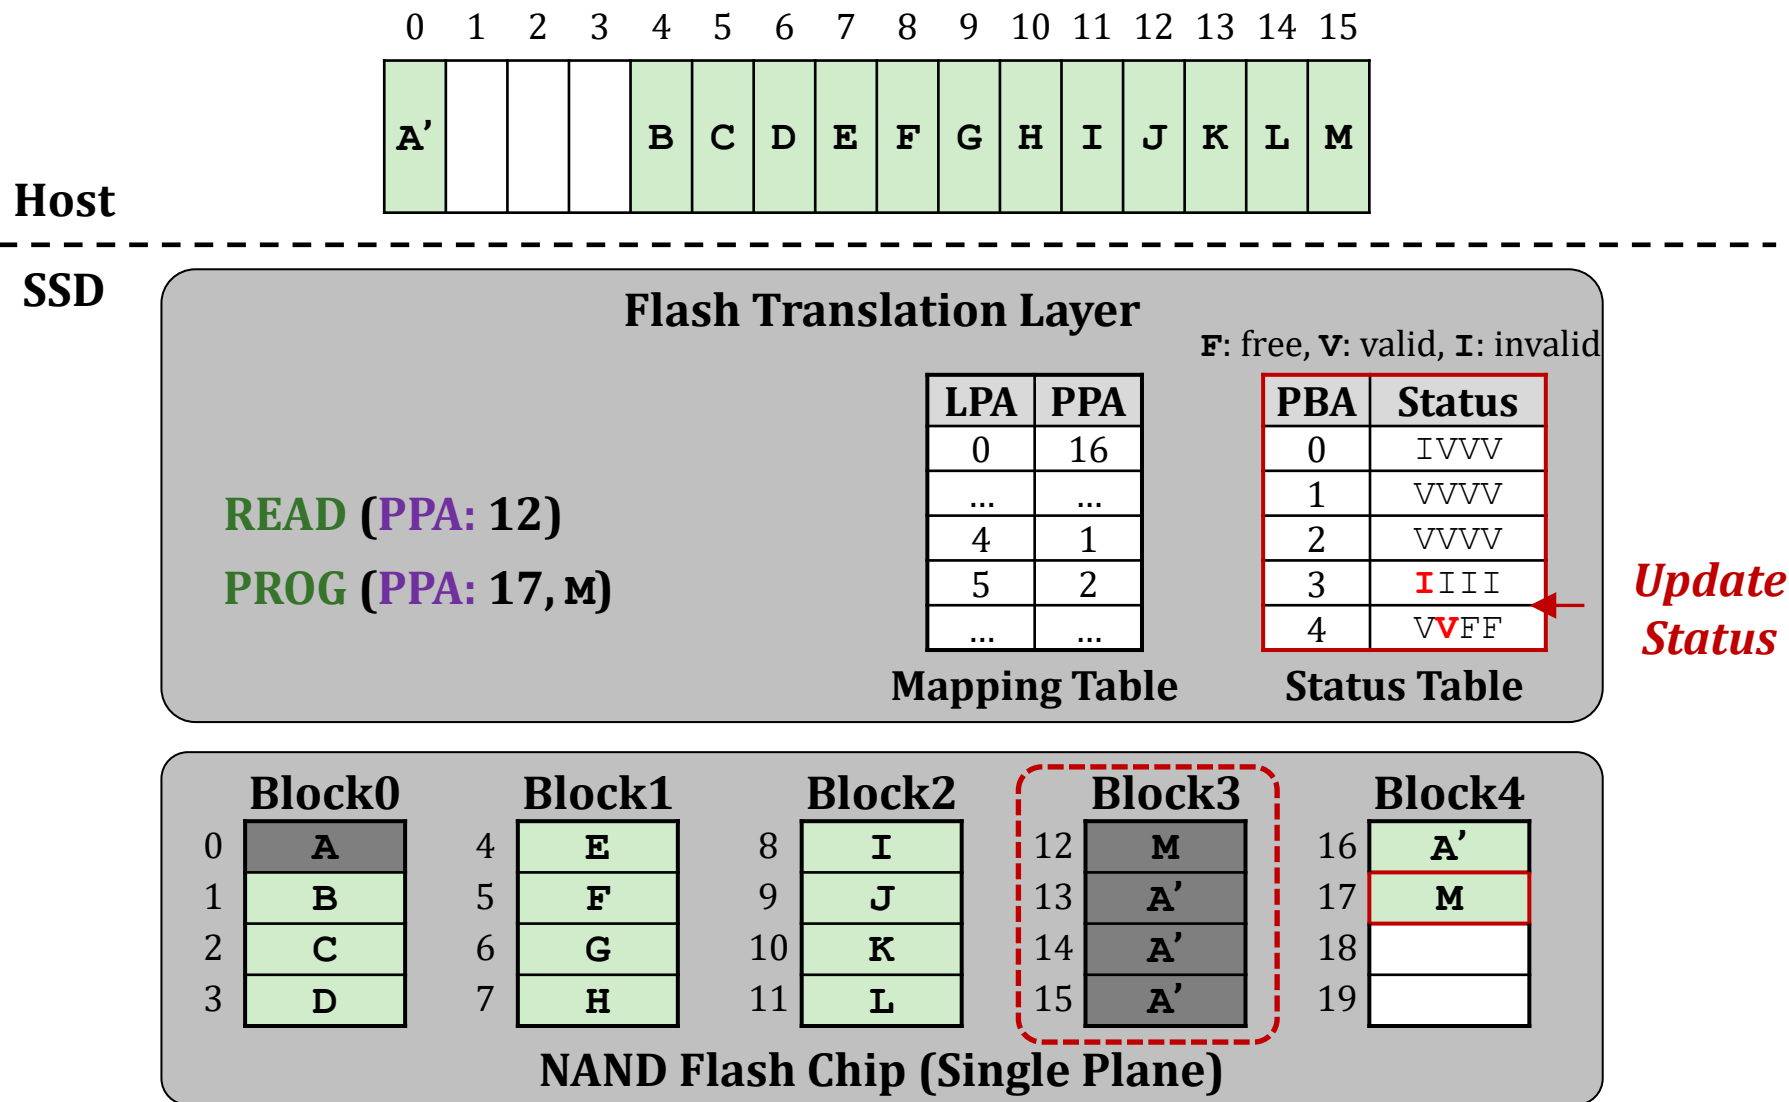
Write Request Handling: Garbage Collection



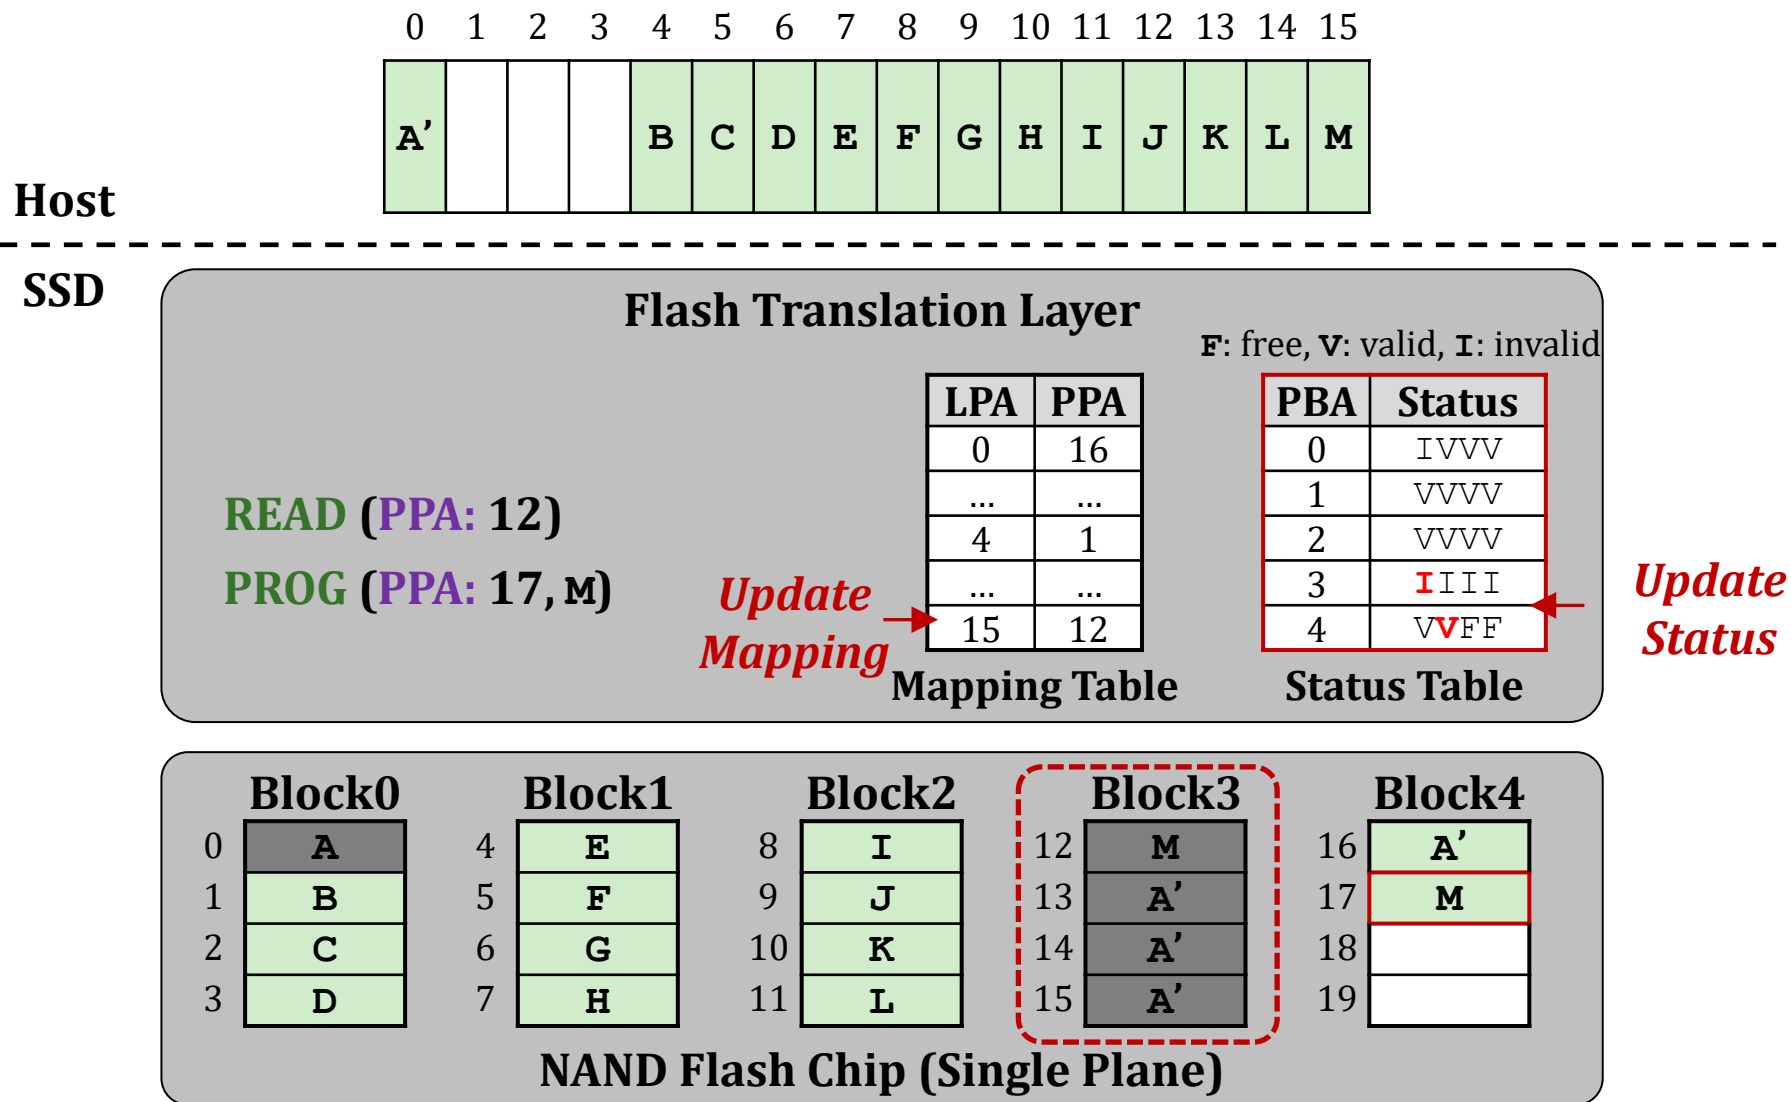
Write Request Handling: Garbage Collection



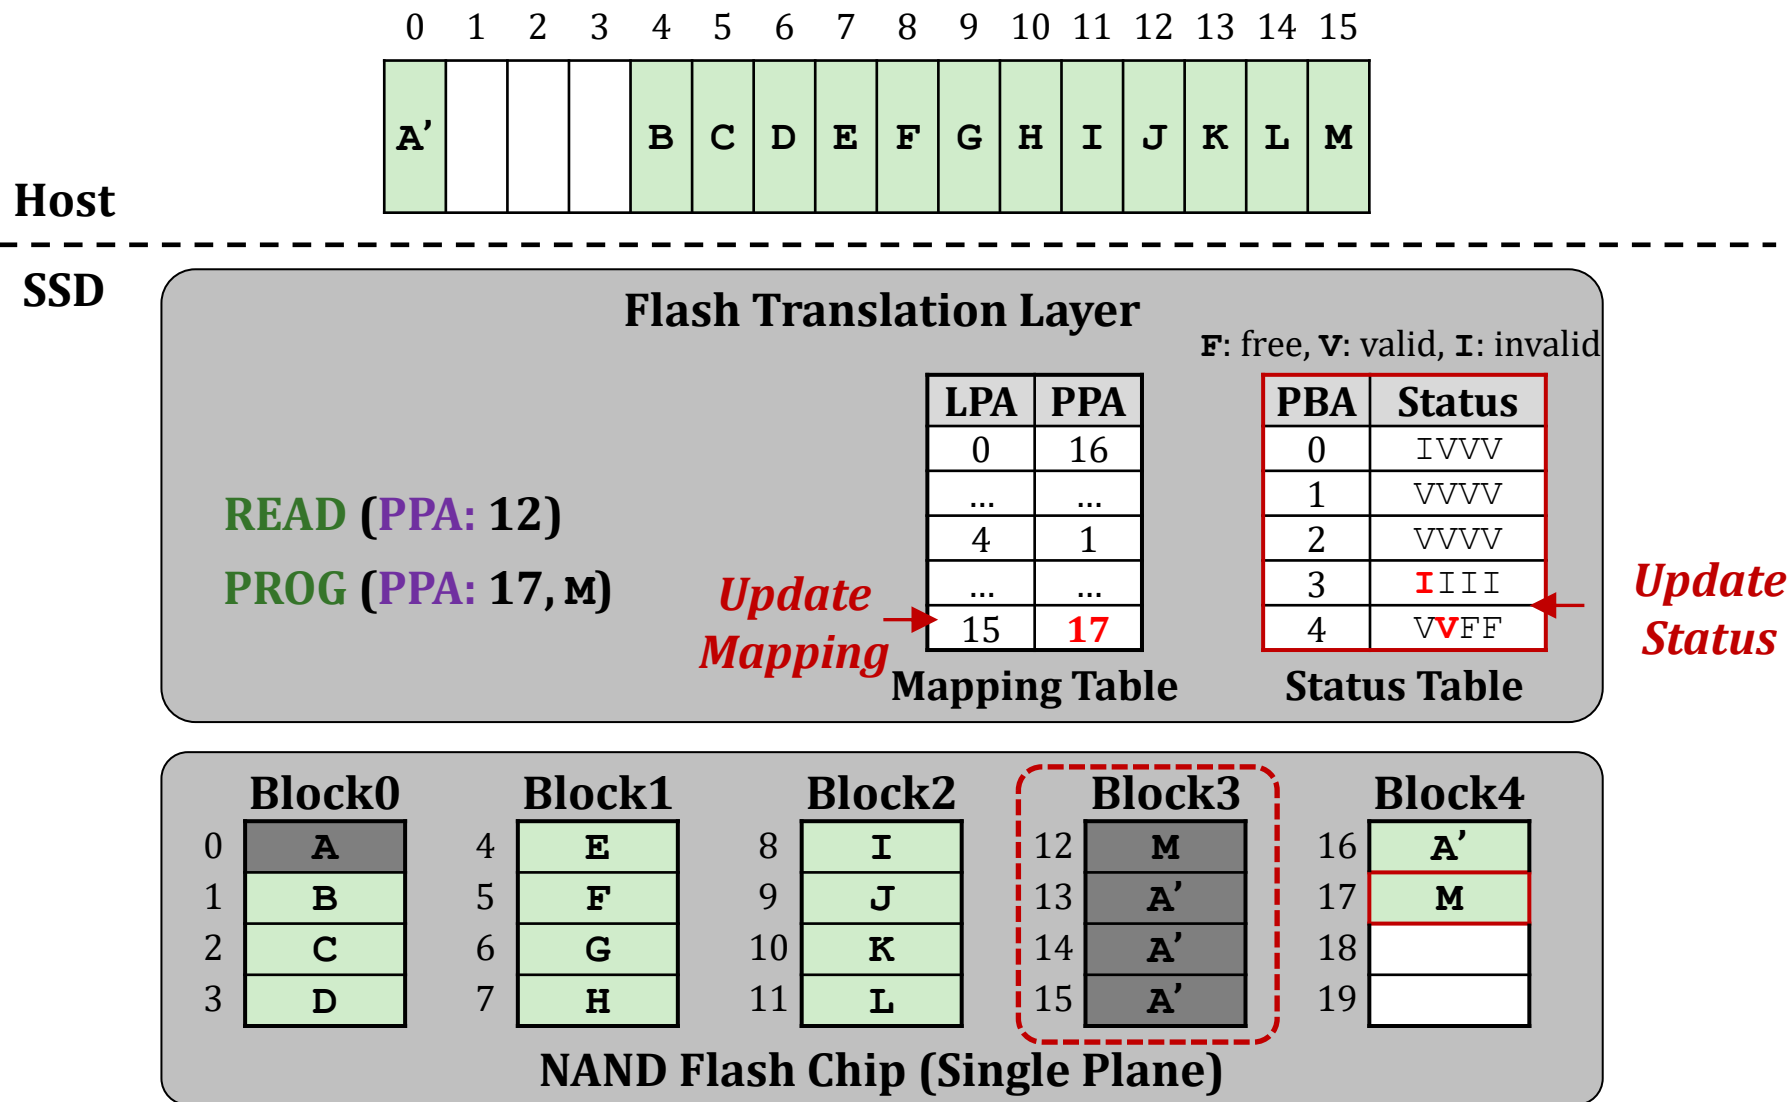
Write Request Handling: Garbage Collection



Write Request Handling: Garbage Collection



Write Request Handling: Garbage Collection



Write Request Handling: Garbage Collection

- **Q:** How FTL knows PPA 12 (data **M**) is mapped to LPA 15?
 - Unless it maintains **P2L mappings**?
- **A:** P2L mapping is stored in each physical page's **OOB (Out-of-Band)** area

READ (PPA: 12)

PROG (PPA: 17, M)

Update Mapping

Mapping Table


LPA	PPA
0	16
...	...
4	1
...	...
15	17

F: free, V: valid, I: invalid

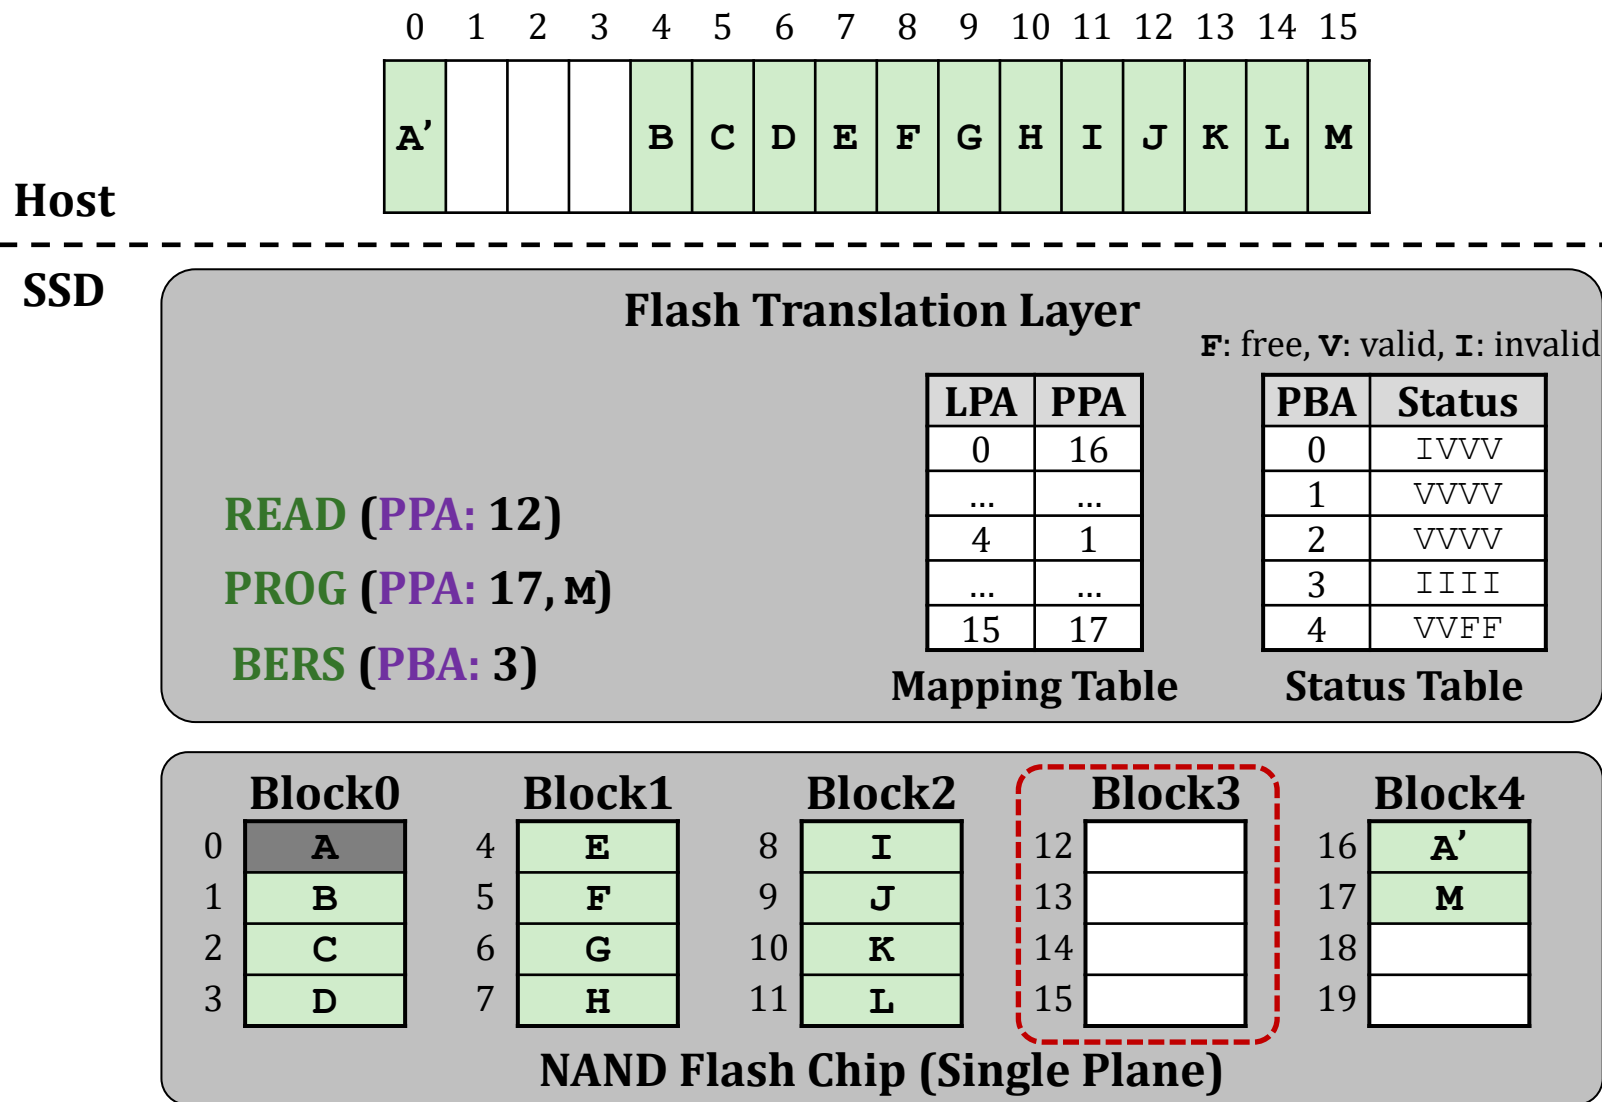
PBA	Status
0	I V V V
1	V V V V
2	V V V V
3	I I I I
4	V V F F

Status Table

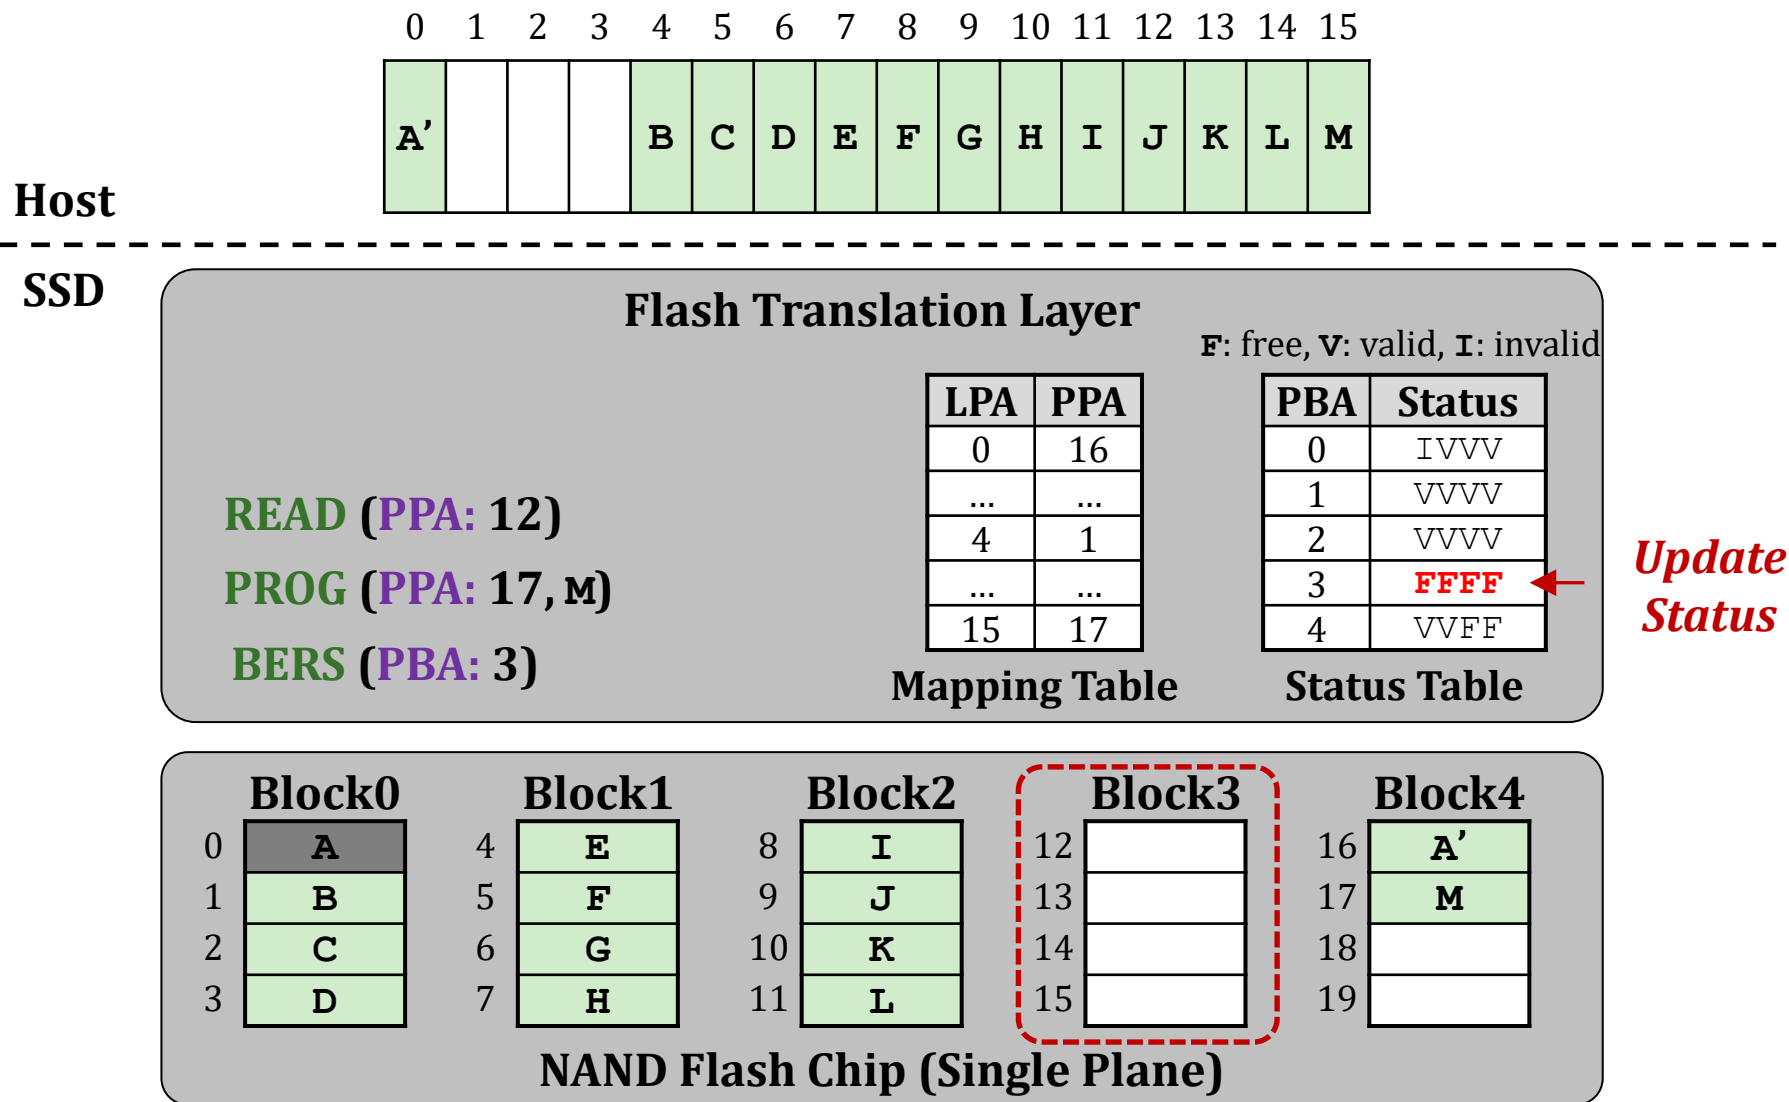
Update Status



Write Request Handling: Garbage Collection



Write Request Handling: Garbage Collection



Write Request Handling: Garbage Collection

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15

Note:

- Block erasure (and status update) is done **just before** programming a new page to the block (i.e., **lazy erase**)
 - Due to the **open-block** problem

(PPA: 12)

(PPA: 17, M)

(PBA: 3)

Mapping Layer

F: free, V: valid, I: invalid

LPA	PPA
0	16
...	...
4	1
...	...
15	17

Mapping Table

PBA	Status
0	IVVV
1	VVVV
2	VVVV
3	FFFF
4	VVFF

Status Table

Update Status

Block0

0	A
1	B
2	C
3	D

Block1

4	E
5	F
6	G
7	H

Block2

8	I
9	J
10	K
11	L

Block3

12	
13	
14	
15	

Block4

16	A'
17	M
18	
19	

NAND Flash Chip (Single Plane)

Performance Issues

- Garbage collection **significantly affects** SSD performance
 - High latency: **Large block size** of modern NAND flash memory
 - Assume 1) a block contains **576** pages,
2) only **5%** of the pages in the victim block are valid
3) $t_R = 100\text{ us}$, $t_{\text{PROG}} = 700\text{ us}$, $t_{\text{BERS}} = 5\text{ ms}$
 - # of pages to copy = $576 \times 0.05 = 28.8 \rightarrow 28$ pages
 - GC latency $> 28 \times (t_R + t_{\text{PROG}}) + t_{\text{BERS}} = \mathbf{27,400\text{ us}}$
 - **Order(s) of magnitude larger** latency than t_R and t_{PROG}
 - Copy operations are **the major contributor** (rather than t_{BERS})
 - If FTL performs GC in an **atomic** manner,
it **delays** user requests for a **significantly long time**
 - Long **tail latency** (performance fluctuation)
 - **Noisy neighbor**: a read-dominant workload's performance would be significantly affected when running with a write-intensive workload (+ performance fairness problem)

Performance Issues: Mitigation

- **TRIM** (**UNMAP** or **discard**) command
 - ❑ Informs FTL of **deletion/deallocation** of a logical block
 - ❑ Allows FTL to **skip copy** of obsolete (i.e., invalid) data
- **Background GC**: Exploits SSD idle time
 - ❑ Challenge: how to accurately predict SSD idle time
 - ❑ Premature GC: copied pages could have been invalidated by the host system
- **Progressive GC**: Divide GC process into subtasks
 - ❑ e.g., copying 28 pages \rightarrow (copying 1 page + servicing user request) \times 28
 - ❑ Effective at decreasing tail latency

Required Materials

- Address Mapping

- Aayush Gupta, Yongjae Kim, and Bhuvan Urgaonkar, “DFTL: A Flash Translation Layer Employing Demand-based Selective Caching of Page-level Address Mappings,” In ASPLOS 2009.

Recommend Materials

- Cache read & Read-retry
 - Jisung Park, Myungsuk Kim, Lois Orosa, Jihong Kim, and Onur Mutlu, "[Reducing Solid-State Drive Read Latency by Optimizing Read-Retry](#)," In ASPLOS 2021.

- Program & Erase Suspension
 - Guanying Wu and Xunbin He, "[Reducing SSD Read Latency via NAND Flash Program and Erase Suspension](#)," In USENIX FAST 2012.
 - Shine Kim, Jonghyun Bae, Hakbeom Jand, Wenjing Jin, Jeonghun Gong, Seungyeon Lee, Tae Jun Ham, and Jae W. Lee, "[Practical Erase Suspension for Modern Low-latency SSDs](#)," In USENIX ATC 2019.

P&S Modern SSDs

Advanced NAND Flash Commands & Address Translation

Dr. Jisung Park

Prof. Onur Mutlu

ETH Zürich

Spring 2022

4 April 2022